

FLAMA: A PYTHON FRAMEWORK FOR DEVELOPMENT AND DEPLOYMENT OF PRODUCTION-READY APIs, MACHINE LEARNING, AND LLM SERVICES

José A. Perdigüero López, Miguel A. Durán-Olivencia

Vortico Tech, Málaga, Spain

We present FLAMA, an open-source Python framework designed for the development and deployment of production-ready web APIs, machine learning services, and large-language-model (LLM) applications. Built on the Asynchronous Server Gateway Interface (ASGI) standard, FLAMA provides a type-driven, async-first programming model that unifies traditional REST API development, predictive model serving, and generative AI inference within a single, coherent architecture.

The framework is organized around seven core subsystems. A component-based dependency injection system resolves handler parameters from type annotations at startup, replacing ad-hoc patterns (global singletons, request-attached state, middleware side effects) with a uniform, testable mechanism. A pluggable schema layer supports three validation libraries (Pydantic, Marshmallow, and Typesystem) through an adapter that normalizes field introspection, validation, and JSON Schema emission. An automatic CRUD resource generator transforms a SQLAlchemy table definition and a schema class into a complete set of REST endpoints, backed by the Repository and Unit of Work domain-driven design patterns for transactional consistency. A portable binary format (.flm) packages trained models from scikit-learn, TensorFlow, PyTorch, and Hugging Face Transformers together with compressed metadata, hyperparameters, training metrics, and auxiliary artifacts into self-describing files that can be deployed with zero application code. A multi-backend LLM serving system deploys generative models through vLLM (Linux/CUDA) or MLX (Apple Silicon), exposing them simultaneously under four wire-protocol dialects (OpenAI, Anthropic, Ollama, and a native streaming protocol) with a shared codec and decoder pipeline that normalizes reasoning channels and tool-call extraction across heterogeneous model families. A Rust-accelerated core, compiled via Maturin, provides high-throughput route resolution, path and host matching, JSON encoding, compression, and cookie and multipart parsing as native Python extensions. Finally, a Model Context Protocol (MCP) module enables any FLAMA application to act as an MCP server, exposing tools, resources, and prompts to AI model clients over JSON-RPC 2.0.

Built-in capabilities include JWT authentication with permission-based middleware, two pagination strategies (page-number and limit-offset), background task execution in threads or processes, WebSocket endpoints with encoding negotiation, Server-Sent Event and Newline-Delimited JSON streaming responses, OpenAPI 3.2.0 schema generation from handler signatures, and a command-line interface for running applications, serving models, downloading and packaging models from remote repositories, performing offline inspection and inference, and migrating codebases across major versions.

We describe the architecture in detail, present the programming model through worked examples, and compare FLAMA with existing frameworks, model serving platforms, and LLM inference engines. This document serves as both a comprehensive technical reference and an extended tutorial for practitioners seeking to build and deploy ML-powered and LLM-powered APIs with minimal operational overhead.

Address: Vortico Tech, Málaga, 29100, Spain

Date: August 2026

Correspondence: research@vortico.tech



Part I

Foundations

1 Introduction

1.1 The deployment gap in machine learning

Fitting a model is no longer the hard part. A decade of work on training frameworks, pre-trained model repositories, and experiment-tracking platforms has seen to that: between scikit-learn ([Pedregosa et al., 2011](#)), TensorFlow ([Abadi et al., 2016](#)), PyTorch ([Paszke et al., 2019](#)), and the Hugging Face ecosystem ([Wolf et al., 2020](#)), a competent practitioner can have a working classifier in an afternoon.

Putting that model in front of users is a different matter. [Sculley et al. \(2015\)](#) made the point with a diagram that has been reproduced ever since: a small box marked *ML code*, surrounded by much larger ones for data collection, feature extraction, configuration, serving infrastructure, and monitoring. Those proportions have not improved. Surveying case studies seven years later, [Paleyes et al. \(2022\)](#) still found deployment and integration consuming more of a project than the modelling did.

The cause of all this is architectural. Researchers' tools and engineers' tools grew up apart, in different communities, assuming different things about the runtime, the interface contract, and the operational lifecycle. What that separation produces, in practice, is a fragmented toolchain, and a typical production ML service might combine:

- A *web framework* for HTTP routing, request parsing, input validation, authentication, and error handling (Flask ([Ronacher, 2010](#)), Django ([Django Software Foundation, 2005](#))).
- A *model serving system* for loading trained artefacts, managing model versions, and running inference (TensorFlow Serving ([Google, 2016](#)), TorchServe ([PyTorch Team, 2020](#))).
- An *experiment and packaging layer* for tracking training runs, storing model metadata, and producing deployable artefacts (MLflow ([Databricks, 2018](#)), BentoML ([BentoML Team, 2019](#))).

Every one of those brings its own abstractions, configuration formats, deployment conventions, and operational overhead. The seams between them are where the technical debt accrues: adapter code has to be written and then maintained to translate the web framework's request format into the model server's input format; authentication and authorization get duplicated, or proxied; error handling and logging have to be reconciled across systems that were never designed to cooperate.

1.2 Unifying web APIs and model serving

FLAMA starts from a refusal: web API development and model serving are not distinct concerns, and do not need distinct systems. The motivating observation behind that is a modest one, namely that a production ML service needs what any web API needs plus a short list of model-specific things. Routing, input validation, authentication, pagination, database access, background tasks,

API documentation, and then, on top of all of it, the ability to load a trained model, run inference, return predictions, and inspect metadata. No architectural argument requires those two sets to sit in separate processes, behind separate middleware stacks, or on separate deployment pipelines.

FLAMA is an open-source Python framework offering a single programming model for both. Six ideas organise it, each aimed at a specific pain point in the landscape above:

1. **Type-driven validation.** Request and response schemas are derived from Python type annotations (van Rossum et al., 2014). The framework supports three schema libraries (Pydantic (Colvin, 2017), Marshmallow (Loria, 2013), and Typesystem (Christie, 2019)) through a pluggable adapter layer, so teams are not locked into a single validation ecosystem.
2. **Component-based dependency injection.** All handler dependencies, from database connections and parsed authentication tokens to validated request bodies and model instances, are resolved through a compile-time dependency graph with request-scoped caching. This replaces ad-hoc patterns (global singletons, request-attached state, middleware side effects) with a uniform, testable, and composable mechanism.
3. **Resource abstraction with domain-driven design.** CRUD endpoints over relational tables are generated automatically from a SQLAlchemy (Bayer, 2006) table definition and a schema class, using the Repository and Unit of Work patterns (Evans, 2003; Fowler, 2002). Individual operations can be overridden or extended without abandoning the generated scaffold.
4. **Native predictive model serving.** Trained models from scikit-learn, TensorFlow, PyTorch, and Hugging Face Transformers are packaged into a portable binary format (`.flm`), loaded as injectable components, and exposed through auto-generated inference endpoints. Registering one is a single `add_model()` call, and the artifact’s own metadata decides which backend loads it.
5. **First-class LLM serving.** Large language models are deployed through a multi-backend engine (vLLM on Linux/CUDA, MLX on Apple Silicon) and exposed simultaneously under four wire-protocol dialects (OpenAI, Anthropic, Ollama, and a native streaming protocol) through a shared codec and decoder pipeline that normalizes reasoning channels, tool-call extraction, and multimodal content across heterogeneous model families. Each dialect mounts under its own prefix, so a single model answers at `/openai/v1/chat/completions` and `/anthropic/v1/messages` at once.
6. **Production-ready defaults.** JWT authentication, pagination, background task execution in threads or processes, OpenAPI 3.2.0 schema generation, debug tooling, and a six-command CLI are included without additional dependencies or configuration.

1.3 Scope and audience

This document does two jobs at once. It is a technical reference, covering architecture, programming model, and implementation in enough depth for contributors and advanced users, and it is also an extended tutorial, walking each subsystem through worked examples a practitioner can lift into their own project. Those two jobs pull in slightly different directions, and where they conflict we have favoured the reference.

Three kinds of reader are in view:

- *ML engineers* who need to deploy trained models as production services without learning a separate serving infrastructure.

- *Backend engineers* who build REST APIs and want built-in support for data validation, CRUD generation, domain-driven design, and dependency injection.
- *Framework designers* interested in the architectural decisions behind a unified API-and-ML framework, including the pluggable schema adapter, the component-based DI system, and the Rust-accelerated routing core.

1.4 Document structure

The document is organized into five parts:

Part I: Foundations introduces the framework, states the design principles, and describes the layered architecture, the request processing pipeline, and the module system (Sections 1–3).

Part II: The web framework covers the core API-building subsystems: routing and endpoints, schema and data validation, dependency injection, resources with domain-driven design, authentication and authorization, pagination, background tasks, and middleware (Sections 4–11).

Part III: Machine learning and generative AI presents the predictive model serving subsystem (the `.flm` serialization format, framework-specific model wrappers, three levels of serving abstraction), the generative AI stack (multi-backend LLM serving, wire-protocol dialects, the codec and decoder pipeline, streaming), and the Model Context Protocol (Section 12).

Part IV: Operations and tooling describes the FLAMA CLI in detail (including model acquisition, codebase migration, and LLM serving commands), the configuration management system, deployment patterns, debug mode, and the testing infrastructure (Sections 19–22).

Part V: Ecosystem and outlook positions FLAMA relative to existing frameworks and LLM inference engines through a detailed feature comparison, and outlines the roadmap for future development (Sections 24–25).

The framework is released under the Apache 2.0 licence and is available at <https://github.com/vortico/flama>, with documentation at <https://flama.dev> and package distribution via PyPI (`pip install flama`).

2 Design principles

A small set of principles governs every architectural decision in FLAMA. Because they constrain the design space, and because they are what separates the framework from the alternatives that made different trades, it is worth setting them out explicitly rather than leaving them to be reconstructed from the code. None of them is abstract. Each came out of a concrete problem met in real API and ML deployment work.

2.1 Asynchronous by default

FLAMA targets the ASGI (Asynchronous Server Gateway Interface) specification (Godwin and ASGI contributors, 2016), which models an application as an asynchronous callable:

```
1 async def application(scope: Scope,
2                       receive: Receive,
3                       send: Send) -> None:
```

4 . . .

Three arguments arrive. The scope is a dictionary describing the connection, giving its type, path, headers, and query string, while the other two are channels: `receive` is where request body chunks come from, and `send` is where response headers and body chunks go. The whole interface is the asynchronous answer to WSGI, which has carried the Python web ecosystem since PEP 333 (Eby, 2003) and is synchronous by construction, since a WSGI application holds its thread for the whole of each request.

Everything inside FLAMA is written against `async/await`, from middleware dispatch down through database queries to model inference. Two consequences follow, one for performance and one for design:

1. **I/O multiplexing.** I/O-bound workloads (database queries, HTTP calls to upstream services, file reads, network waits) can be multiplexed on a single thread without blocking. A single worker process can serve thousands of concurrent connections, because each connection yields control to the event loop during I/O waits rather than holding a thread hostage.
2. **Explicit concurrency boundaries.** CPU-bound work (model inference, data transformation, report generation) can be offloaded to background threads or processes through the framework's concurrency primitives (`BackgroundThreadTask`, `BackgroundProcessTask`) without breaking the async contract. The dispatch itself is handled by `flama.concurrency.run()`, which awaits a coroutine directly and sends anything else to a thread pool with the caller's `contextvars` context copied across.

None of which forces async syntax on anyone. A synchronous handler is wrapped in `asyncio.to_thread()` at resolution time, so an ordinary function works as well as a coroutine:

```
1 @app.route("/health")
2 def health_check() -> dict:
3     return {"status": "ok"}
```

`health_check` is not a coroutine, the framework notices as much, and it gets scheduled on a thread pool. The practical payoff is migration: porting an existing synchronous codebase does not mean rewriting every handler before it will run.

2.2 Type annotations as the source of truth

Python type annotations (van Rossum et al., 2014; Gonzalez et al., 2016; Langa, 2019; Prados and Moss, 2019) serve as the single source of truth for four distinct concerns in FLAMA:

1. **Request parsing.** The names and types of handler parameters determine where each value is extracted from (path segment, query string, or request body) and how it is parsed.
2. **Validation.** Schema classes referenced in type annotations trigger automatic validation of the corresponding request data. Validation errors produce structured 400 Bad Request responses.
3. **Dependency resolution.** Parameters whose types match a registered component are resolved through the dependency injection system. The type annotation is the sole mechanism by which a handler declares its dependencies.

4. **API documentation.** The OpenAPI 3.2.0 specification is generated from handler signatures at startup: parameter types become OpenAPI parameter schemas, return types become response schemas, and schema classes become reusable component definitions.

The four converge in a way that is easier to show than to describe. Take this handler:

```
1 import typing as t
2 from flama import schemas
3
4 @app.route("/users/{user_id:int}", methods=["PUT"])
5 async def update_user(
6     user_id: int,
7     data: t.Annotated[schemas.Schema, schemas.SchemaMetadata(UserUpdate
8         )],
9     connection: AsyncConnection,
10    token: AccessToken,
11) -> t.Annotated[schemas.Schema, schemas.SchemaMetadata(User)]:
12    ...
```

From this signature alone, the framework determines that:

- `user_id` is an integer extracted from the URL path segment `{user_id:int}`.
- `data` is validated from the JSON request body against the `UserUpdate` schema, using the configured schema library; the `SchemaMetadata` annotation is what tells the framework which parameter carries the request body and which schema to validate it against.
- `connection` is an `AsyncConnection` resolved by a registered component, typically one backed by the `SQLAlchemyModule`.
- `token` is an `AccessToken`, decoded by the JWT authentication component from the request's `access_token` header or cookie.
- The return annotation fixes the response schema (`User`) and, with it, the OpenAPI response definition.

Nothing else is required beyond the route binding itself: no decorators, no configuration dictionaries, no registration calls. The signature alone is the endpoint's contract.

2.3 Pluggable schema libraries

Picking a validation library on a user's behalf is a decision that ages badly. FLAMA declines to make it, and interposes an adapter over three established ones instead, each available as an optional install extra:

- **Pydantic** (Colvin, 2017): the most widely adopted validation library in the Python ecosystem, with Rust-accelerated core validation.
- **Marshmallow** (Loria, 2013): a mature library with a rich ecosystem of plugins and a serialization API oriented around explicit field declarations.
- **Typesystem** (Christie, 2019): a lightweight library designed for data validation and form rendering in web frameworks.

Whichever is chosen, the adapter translates its field and schema representations into one internal model, the `Field` and `Schema` structures of Section 5, and everything downstream (validation, JSON Schema emission, OpenAPI generation) reads that instead. The choice is made once, at application initialization:

```
1 app = Flama(schema_library="pydantic")
```

after which the library's own classes are used throughout the codebase and the translation stays out of sight. Two things are bought with this arrangement: a team already invested in one library can adopt FLAMA without rewriting its schemas, and the core validation pipeline stops being hostage to any single library's API churn. The second is not hypothetical; Pydantic's 1.x-to-2.x transition rewrote most of its public surface, and an adapter is where a framework absorbs that kind of change instead of passing it on.

2.4 Dependency injection as a first-class mechanism

How does a handler get hold of what it needs? Python web frameworks have accumulated several answers, none of them uniform: module-level singletons, state hung off the request (`request.state.db` and relatives), middleware that quietly writes into the ASGI scope, explicit `Depends()` calls that weld the handler to one resolution strategy.

In place of all of them FLAMA uses component-based dependency injection, along the lines [Fowler \(2004\)](#) sets out. A handler has no business knowing *how* its dependencies get built; what it does instead is declare, through type annotations, *what* it wants, leaving the framework to supply the instances at call time.

Three things follow:

1. **Testability.** Swapping a dependency in a test means registering a different component, with no patching and no monkeypatching involved. A test wanting a fake database registers a component that hands back a mock connection, and the handler is none the wiser.
2. **Composability.** Components can depend on other components, so `UserRepository` can ask for an `AsyncConnection` that is itself produced by a connection component, and the framework works out the resulting graph on its own.
3. **Uniformity.** One mechanism covers database connections, validated request data, authentication tokens, model instances, pagination parameters, and anything a user defines. One resolution system, one caching strategy, one place to look when a dependency misbehaves.

2.5 Convention over configuration for CRUD

CRUD endpoints over a relational table are the most written and least interesting code in API development. Most frameworks ask for five to seven handler functions, each with a route binding, input validation, error handling, and its own database interaction. Almost all of it is boilerplate, since the handlers take their shape from the table, and departures from the standard pattern are rare enough to be worth handling as exceptions rather than designing for.

So FLAMA generates them, from two declarations: a SQLAlchemy ([Bayer, 2006](#)) table and a schema class. Create, retrieve, update, delete, list, and bulk-drop all come out with the right status codes, input validation, pagination, and error handling (integrity errors become 400 Bad Request, missing

resources 404 Not Found). The generated code is built on the domain-driven design layer of Section 7, and any single operation can be overridden or extended without giving up the rest.

In the usual case, hundreds of lines of handler code collapse into a class declaration of under ten.

2.6 ML and LLM serving as native concerns

Model serving in FLAMA isn't a plugin, a sidecar, or an afterthought. It's a module, on the same footing as routing, schema generation, and resource management. Predictive models and large language models get the same treatment:

- A trained model is packaged into a `.flm` file (a compressed binary containing the model artefact, framework metadata, training parameters, evaluation metrics, and optional auxiliary artefacts). For predictive models the payload is the serialized weight tensor; for LLMs it is a tarball of the model's checkpoint directory.
- Registration reads only the metadata header, so it stays cheap. Deserialising the body is deferred to a startup event, which lets the server bind its port before a multi-gigabyte checkpoint begins loading. The result is a typed component, injected into handlers like any other.
- A predictive model gets `GET /` for its metadata and `POST /predict/` for inference, both generated. An LLM gets a whole multi-dialect surface instead, following whichever of OpenAI, Anthropic, Ollama, and the native protocol were asked for.
- The same CLI that starts a development server will also serve a model with no application code at all.

A team whose API happens to include a model, then, has no second serving system to operate. The model sits in the same process as everything else, behind the same authentication, under the same logging and monitoring, shipped down the same pipeline.

2.7 Protocol-agnostic generative serving

Generative AI has produced a thicket of mutually incompatible wire protocols: OpenAI's Chat Completions API ([OpenAI, 2023](#)), Anthropic's Messages API ([Anthropic, 2024](#)), Ollama's local inference protocol ([Ollama, 2023](#)), and a long tail of vendor-specific endpoints. Wire a framework's internals to any one of them and its users inherit that choice, along with the integration rewrite that follows every protocol revision.

The way out is to keep the model's execution and the wire format apart. At the centre sits a canonical transport model: a request is a sequence of typed messages carrying multimodal content parts, a response is a stream of typed events (start, text, tool call, trace, stop). A *dialect* stands between that model and the client, and it is three things at once, a parser turning wire format into canonical messages, a renderer turning canonical events into streaming wire frames, and an assembler turning them into buffered response envelopes.

With that in place, one physical model instance serves every dialect at once, each mounted under its own URL prefix, so that the OpenAI SDK, the Anthropic SDK, the Ollama CLI, and a browser pointed at the native chatbot all reach the same application and the same weights. Normalising reasoning channels and pulling tool calls out of whatever shape a model emits them in belongs to the codec and decoder pipeline. That work happens once, upstream of whichever dialect renders the result.

3 Architecture

Three things make up the structure of FLAMA, and each is taken in turn below: the layers it is built from, the path an HTTP request takes down through them, and the module system by which it is extended.

3.1 Layered design

The framework is layered, with every layer handing the one above it a set of abstractions, and every layer able to be reasoned about on its own, which is rather the point of the arrangement. Figure 1 shows the whole of it.




Figure 1 Layered architecture of FLAMA. The ASGI protocol provides the server interface. The Flama Core layer, with Rust-compiled performance-critical paths, supplies HTTP abstractions, route resolution, path and host matching, JSON encoding, compression, and multipart and cookie parsing. The Application layer adds dependency injection, component management, and lifecycle hooks. Modules extend the application with schema generation, CRUD resources, ML and LLM model serving, and the Model Context Protocol.

The layers, from bottom to top, are:

ASGI Protocol. The interface between an HTTP server, Uvicorn (Christie, 2017) for instance, and the application. FLAMA implements the application side of the contract and no more, and is not itself a server. Any compliant one will host it.

Flama Core. The layer that provides the fundamental HTTP and WebSocket abstractions: **Request**, **Response**, **WebSocket**, middleware composition, static file serving, and content negotiation. Performance-critical paths in this layer are implemented in Rust and compiled as Python extensions via Maturin (PyO3 Project, 2019). The Rust core comprises seven modules: *route resolution* (iterates the registered route entries in a single pass, extracting path parameters as it matches), *path and host matching* (the segment-based template matcher that route resolution calls, plus a host matcher for exact, wildcard-subdomain, and match-any virtual-host patterns), *JSON encoding* (fast serialization with support for datetimes, UUIDs, decimals, dataclasses, and framework types), *compression* (bz2, lzma, zlib, zstd, gzip, and brotli codecs with a streaming compressor), *multipart parsing* (async form-data parsing via Tokio), *cookie handling* (RFC-compliant parsing and Set-Cookie construction), and *HTTP utilities* (content-type parsing). What crosses back into Python is always a plain object, bytes or a tuple or a list or a dictionary,

which keeps FFI overhead down while the inner loops still run compiled. None of this is visible at install time: wheels are published for every supported Python (3.10 through 3.14) on Linux, macOS, and Windows, so no Rust toolchain is needed to `pip install flama`.

Application Layer. The `Flama` class, the object users actually instantiate. It composes a `Router`, an `Injector`, a `Components` registry, a `MiddlewareStack`, an `Events` manager, and a set of `Modules`. Orchestrating the request lifecycle is its job: take the ASGI connection, push it through the middleware stack to the router, resolve dependencies, run the handler, build the response.

Modules. Self-contained units of related functionality. A module can register components, add routes, and hook startup and shutdown. Four ship with the framework (`Schema`, `Resources`, `Models`, `MCP`), and custom ones subclass `Module`.

3.2 The application object

Everything hangs off the `Flama` class, which is itself the ASGI callable, and a single line of it is already a working application:

```
1 from flama import Flama
2
3 app = Flama()
```

That already has OpenAPI schema generation, debug error pages, and an empty route table waiting for handlers. Optional constructor parameters take over from there:

```
1 app = Flama(
2     openapi={
3         "info": {
4             "title": "My API",
5             "version": "1.0.0",
6             "description": "A production API",
7         }
8     },
9     schema_library="pydantic",
10    routes=[...],
11    components=[...],
12    modules=[...],
13    middleware=[...],
14    events={"startup": [...], "shutdown": [...]},
15    debug=False,
16 )
```

Inside, the object is an assembly of six subsystems:

Subsystem	Responsibility
Router	Maps URL paths to handler callables using the Rust-accelerated route table. Manages route registration, path parameter extraction, and method-not-allowed responses.
Injector	Builds and caches dependency resolution trees for each handler at startup. At request time, walks the pre-compiled tree to resolve all handler parameters.
Components	A registry of all component instances available for injection. Includes built-in components (request data, validation, authentication) and user-registered components.
MiddlewareStack	An ordered chain of ASGI middleware that wraps the router. Processes requests in registration order and responses in reverse order.
Modules	A dictionary of named module instances (Schema, Resources, Models, MCP). Each module is accessible as an attribute of the application (e.g. <code>app.models</code>).
Events	Startup and shutdown event handlers, executed when the ASGI lifespan protocol signals that the application should initialize or tear down.

3.3 Request processing pipeline

A request moves through FLAMA in a fixed sequence of stages, and the sequence is worth knowing, because middleware ordering, the timing of dependency resolution, and the point at which errors surface all follow from it. Figure 2 traces the flow.




Figure 2 Request processing pipeline. The request flows left to right through the middleware stack, route resolution, dependency injection, and handler execution. The response flows back through the middleware stack before being sent to the client.

There are eight stages, described below in the order they run.

3.3.1 Stage 1: ASGI reception

The server, Uvicorn or another, accepts the connection and builds a scope dictionary out of it: method, path, headers, query string, client address. It calls the application’s `__call__` with that scope and the `receive/send` channels.

3.3.2 Stage 2: Middleware processing (request phase)

The middleware stack is an ordered sequence of ASGI callables, each wrapping the next, and the outermost sees the request first. Any of them can inspect or modify the scope, read the body through `receive`, hand control onward, or cut the pipeline short by answering directly, which is what an authentication failure or a CORS preflight does.

3.3.3 Stage 3: Route resolution

The **Router** hands the path to the Rust-compiled `RouteTable` and gets back one of three outcomes:

- **Full match:** the path matches a registered route pattern, and the HTTP method is allowed. The router extracts the path parameters (typed according to the route pattern) and dispatches to the matched handler.

- **Method not allowed:** the path matches a route, but the HTTP method is not supported. The router returns a 405 Method Not Allowed response.
- **No match:** the path does not match any route. The router returns a 404 Not Found response.

3.3.4 Stage 4: Dependency injection

For the matched handler, the `Injector` walks the pre-compiled resolution tree, built at startup as described in Section 6, and every parameter in the signature resolves from one of three sources:

1. **Context values:** objects available directly from the ASGI scope or the request context (`Request`, `Flama`, `Route`, `WebSocket`).
2. **Components:** registered component instances whose `resolve()` method produces a value of the required type. Components may have their own upstream dependencies, which are resolved recursively.
3. **Parameters:** primitive values extracted from path segments or query strings.

Resolved values land in a per-request cache. A component declaring `cacheable = True`, which is the default, resolves at most once per request no matter how many parameters or nested components ask for it.

3.3.5 Stage 5: Handler execution

With the arguments resolved, the handler runs, or, for a class-based endpoint, the method matching the verb runs in its place. Should that handler be synchronous rather than a coroutine, it goes to a thread pool via `asyncio.to_thread()`.

3.3.6 Stage 6: Response building

Whatever the handler returned now becomes an HTTP response. A `Response` object is used as it stands. Anything else is wrapped in an `APIResponse` with the route's declared response schema attached. When that schema is present, `APIResponse` re-validates the value through the schema adapter before encoding it to JSON, so a handler whose return type is schema-annotated must return a plain `dict` (or a list of dicts, for a collection response) with the schema's field names as keys, not an instantiated schema object.

3.3.7 Stage 7: Middleware processing (response phase)

The response travels back out through the same stack in reverse order, and along the way headers get inspected and rewritten, metrics logged, bodies compressed, and cleanup done.

3.3.8 Stage 8: ASGI transmission

Headers and body go out over the `send` channel. Any background tasks attached to the response run once transmission has finished.

3.4 The module system

A module is a self-contained unit of related functionality: a subclass of `Module`, implementing whichever lifecycle hooks it needs and ignoring the rest.

```

1 from flama.modules import Module
2
3 class AnalyticsModule(Module):
4     name = "analytics"
5
6     def __init__(self, dsn: str) -> None:
7         super().__init__()
8         self._dsn = dsn
9         self.client: AnalyticsClient | None = None
10
11     async def on_startup(self) -> None:
12         self.client = AnalyticsClient(self._dsn)
13
14     async def on_shutdown(self) -> None:
15         await self.client.flush()
16         self.client = None

```

A module holds its own state rather than writing it onto the application: the instance is reachable from the application under its `name` (here `app.analytics`), so whatever it exposes as attributes or methods becomes the module's public surface. Modules can also register components and routes dynamically during startup, which makes them a natural unit for packaging a whole subsystem (database connection management, model lifecycle, schema generation) for reuse.

The framework ships with four modules:

SchemaModule generates the OpenAPI 3.2.0 specification ([OpenAPI Initiative, 2025](#)) from route metadata and type annotations. It walks all registered routes at startup, extracts parameter and body annotations, and emits a complete OpenAPI document. The specification is served as JSON at a configurable path (default `/schema/`), and an interactive Swagger UI is rendered at `/docs/`.

ResourcesModule manages the lifecycle of REST resources (Section 7), exposing `add_resource()` to attach a resource class to the application under a path prefix. Database connectivity itself is a separate, optional module (`SQLAlchemyModule`) that a `CRUDResource`-based application registers alongside it.

ModelsModule manages the lifecycle of ML models (Section 12): it registers `ModelComponent` instances for injection, mounts their endpoints, and materialises each one during startup. That last step runs sequentially rather than concurrently, on the grounds that loading several multi-gigabyte artifacts at once trades a slow cold start for an out-of-memory failure.

MCPModule implements the Model Context Protocol (Section 18), enabling a FLAMA application to act as an MCP server that exposes tools, resources, and prompts to AI model clients over a JSON-RPC 2.0 transport.

Custom modules are registered at application construction time as a list of module instances, where each subclass supplies its own `name` attribute, which is how the application indexes it:

```

1 app = Flama(
2     modules=[AnalyticsModule(dsn=ANALYTICS_DSN)],
3 )

```

Once registered, a module is reachable as an attribute of the application, so the example above answers to `app.analytics`. Two modules claiming the same `name` raise at construction rather than silently shadowing one another.

Part II

The Web Framework

4 Routing and endpoints

Routing is the front door of a FLAMA application. It decides which handler sees each incoming request, how path segments become typed parameters, and what happens when nothing matches. Since the route declaration is also where an application's external interface gets pinned down, most of what follows in this part depends on it, which is why it comes first: route types, then the declaration syntax, then the compiled route table underneath, then class-based endpoints.

4.1 Route types

FLAMA distinguishes four kinds of routes, each modelling a different relationship between a URL pattern and a handler:

HTTP routes bind a URL pattern and one or more HTTP methods to a handler callable. The handler may be a plain function (synchronous or asynchronous) or a class-based endpoint (Section 4.5). HTTP routes are the most common route type and account for the majority of endpoints in a typical application.

WebSocket routes bind a URL pattern to a WebSocket endpoint with a three-phase lifecycle: connection, a receive-and-respond loop, and disconnection. WebSocket routes support encoding negotiation (JSON, text, or raw bytes) for incoming messages and can be used for real-time features such as live dashboards and interactive interfaces.

Mounts attach a sub-application or a **Router** instance under a path prefix. Mounts enable modular composition: a team can develop a self-contained set of routes in an isolated router and mount it into the main application at deployment time. Mounts also support mounting any ASGI-compliant application, making it possible to combine FLAMA with other ASGI frameworks in a single process.

Resource routes are generated automatically from a resource class declaration (Section 7). They produce a full set of CRUD endpoints, each backed by domain-driven design patterns. Resource routes are syntactic sugar: they expand into a mount containing standard HTTP routes.

4.2 Declaring routes

Routes can be declared with decorator syntax, which is the idiomatic approach for applications where routes are defined close to their handlers:

```
1 import typing as t
2 from flama import Flama, schemas
3
4 app = Flama()
```



```

5
6 UserResponse = t.Annotated[
7     schemas.Schema, schemas.SchemaMetadata(User)
8 ]
9
10 @app.route("/users/{user_id:int}", methods=["GET"])
11 async def get_user(user_id: int) -> UserResponse:
12     """Retrieve a single user by primary key."""
13     ...

```

Every schema-typed value in a signature, whether it travels in the request body or the response, is declared the same way: `t.Annotated[schemas.Schema, schemas.SchemaMetadata(X)]` (Varoquaux and Kashin, 2019), with `list[schemas.Schema]` in place of `schemas.Schema` when the value is a collection. The annotation carries two pieces of information the framework needs and cannot infer from a bare class reference: that the parameter (or return value) is a schema-validated payload at all, and which schema class validates it. Binding these annotations to a named alias, as with `UserResponse` above, keeps signatures readable when the same schema recurs across handlers:

```

1 UserRequest = t.Annotated[
2     schemas.Schema, schemas.SchemaMetadata(UserInput)
3 ]
4
5 @app.get("/users/{user_id:int}")
6 async def get_user(user_id: int) -> UserResponse:
7     ...
8
9 @app.post("/users/")
10 async def create_user(data: UserRequest) -> UserResponse:
11     ...
12
13 @app.put("/users/{user_id:int}")
14 async def update_user(
15     user_id: int, data: UserRequest
16 ) -> UserResponse:
17     ...
18
19 @app.delete("/users/{user_id:int}")
20 async def delete_user(user_id: int) -> None:
21     ...

```

A handler that returns a `Response` object directly, or that has no documented payload, simply omits the return annotation; the framework then generates no response schema for it.

For applications where routes are defined separately from their handlers (e.g. in a configuration module or a factory function), routes can be constructed programmatically:

```

1 from flama.routing import Route, Router, Mount
2
3 routes = [
4     Route("/users/{user_id:int}", get_user,
5         methods=["GET"]),
6     Route("/users/", create_user,
7         methods=["POST"]),

```

```

8     Mount("/admin", app=admin_router),
9 ]
10
11 app = Flama(routes=routes)

```

4.3 Path parameters and type converters

Path segments enclosed in braces are treated as parameters. Each parameter may include a type converter, separated by a colon, that determines how the raw string segment is parsed and validated:

Pattern	Type	Matching behaviour
{id}	<code>str</code>	Matches any single path segment (no slashes).
{id:int}	<code>int</code>	Matches a sequence of digits; returns an integer.
{id:float}	<code>float</code>	Matches a decimal number; returns a float.
{id:uuid}	<code>UUID</code>	Matches a UUID string (8-4-4-4-12 hex); returns a <code>uuid.UUID</code> object.
{id:decimal}	<code>Decimal</code>	Matches a decimal number; returns a <code>decimal.Decimal</code> object.

Type converters are applied during route resolution, before the handler is called. If a path segment does not match the declared type (e.g. `/users/abc` against `{id:int}`), the route is not matched and the router continues searching for alternative routes or returns 404.

4.4 The Rust-accelerated route table

Route resolution runs on every request, so its cost is paid on every request, and a router that degrades as the table grows degrades the whole application with it. FLAMA delegates the matching to a Rust-compiled `RouteTable`, built with Maturin and exposed to Python as a native extension module. What that buys is measurable: under Callgrind, a full request cycle costs about 3.54 M estimated cycles against a table of ten routes and 3.69 M against a table of two hundred, and within that two-hundred-route table the difference between matching the first entry and the last is under 3%. Table size, in other words, has almost stopped mattering.

The `RouteTable` stores route patterns in a pre-compiled data structure and resolves incoming paths by iterating over all registered entries, performing segment-by-segment matching for each candidate. Constant segments are compared as byte slices, while parameterized segments are validated against their declared type (integer, UUID, etc.). Path parameter extraction is performed in the same pass as pattern matching, avoiding a second traversal of the URL. The table exposes three operations:

- `add_entry(pattern, methods, index)`: registers a route pattern with its allowed HTTP methods and an index that maps back to the Python route object.
- `resolve(path, method)`: matches a request path against all registered patterns and returns a `ResolveResult` containing the match type (full, mount, or not found), the route index, the extracted path parameters, and the set of allowed methods.
- Path parameter extraction is performed in the same pass as pattern matching, avoiding a second traversal of the URL.

The Rust JSON encoder (`json_encoder`) serves the same purpose for response serialization: it replaces Python's `json.dumps()` with a compiled implementation that handles the common case (flat dictionaries of strings and numbers) significantly faster.

4.5 Class-based endpoints

For handlers that serve multiple HTTP methods on the same URL path, or that need to manage a multi-step connection lifecycle, FLAMA provides three base classes: `HTTPEndpoint`, `WebSocketEndpoint`, and `JSONRPCEndpoint`. The first two are the most common and are described below; the third is used internally by the MCP module (Section 18) and follows the same pattern.

4.5.1 HTTP endpoints

An `HTTPEndpoint` groups related handlers into a single class, with one method per HTTP verb:

```
1 import typing as t
2 from flama import schemas
3 from flama.endpoints import HTTPEndpoint
4
5 UserUpdateRequest = t.Annotated[
6     schemas.Schema, schemas.SchemaMetadata(UserUpdate)
7 ]
8
9 class UserEndpoint(HTTPEndpoint):
10     async def get(self, user_id: int) -> UserResponse:
11         """Retrieve a user."""
12         ...
13
14     async def put(
15         self, user_id: int, data: UserUpdateRequest
16     ) -> UserResponse:
17         """Update a user."""
18         ...
19
20     async def delete(self, user_id: int) -> None:
21         """Delete a user."""
22         ...
```

The framework introspects the class at registration time to determine the set of allowed methods. Requests with unsupported methods receive a 405 Method Not Allowed response. HEAD requests are served automatically by any endpoint that implements GET.

Class-based endpoints are registered with the same decorator or programmatic API as function handlers:

```
1 app.route("/users/{user_id:int}")(UserEndpoint)
```

4.5.2 WebSocket endpoints

WebSocket connections (Fette and Melnikov, 2011) follow a three-phase lifecycle modelled by the `WebSocketEndpoint` class:

```
1 from flama.endpoints import WebSocketEndpoint
2
3 class ChatEndpoint(WebSocketEndpoint):
4     encoding = "json"
5
```

```

6     async def on_connect(self, websocket):
7         await websocket.accept()
8
9     async def on_receive(self, websocket, data):
10        response = process_message(data)
11        await websocket.send(json=response)
12
13    async def on_disconnect(self, websocket, websocket_code):
14        cleanup(websocket)

```

The `encoding` attribute determines how incoming messages are parsed ("json", "text", or "bytes"; it is `None` by default, which negotiates). The framework calls `on_connect` when the handshake completes, `on_receive` for each incoming message, and `on_disconnect` when the connection closes. Sending is done through the single `send()` method, which takes the payload as one of the keyword arguments `data`, `json`, or `message` rather than through per-type methods; reading directly, outside the `on_receive` hook, is the mirror-image `receive(data="text")`.

WebSocket routes are registered with the same decorator syntax as HTTP routes:

```

1  @app.websocket_route("/ws/updates")
2  class LiveUpdates(WebSocketEndpoint):
3      encoding = "json"
4
5      async def on_connect(self, websocket):
6          await websocket.accept()
7
8      async def on_receive(self, websocket, data):
9          result = await compute(data)
10         await websocket.send(json=result)
11
12     async def on_disconnect(self, websocket, websocket_code):
13         pass

```

The equivalent programmatic API is available through `app.add_websocket_route(path, endpoint)`. Both approaches produce a `WebSocketRoute` object in the route table.

4.5.3 Streaming responses

For HTTP endpoints that need to produce output incrementally (large file downloads, real-time event streams, or progressive result delivery), FLAMA provides three specialised response classes that stream their content to the client without buffering the entire body in memory.

General streaming. The `StreamingResponse` class accepts a synchronous or asynchronous generator and emits each yielded chunk directly to the client:

```

1  from flama import Flama
2  from flama.http.responses import StreamingResponse
3
4  app = Flama()
5
6  async def generate_report():
7      for chunk in compute_report_chunks():
8          yield chunk

```

```

9
10 @app.get("/report")
11 async def stream_report():
12     return StreamingResponse(
13         generate_report(),
14         media_type="text/csv",
15     )

```

Server-Sent Events (SSE).. The `ServerSentEventResponse` implements the WHATWG Server-Sent Events living standard (WHATWG, 2026), setting the Content-Type to `text/event-stream`, Cache-Control to `no-cache`, and Connection to `keep-alive`. Each yielded item is either a plain string (sent as a bare `data:` line) or a `ServerSentEvent` object, serialised as an SSE frame with optional `id`, `event`, `data`, and `retry` fields. A `ServerSentEvent` with only `comment` set is emitted as a comment-only heartbeat line, keeping the connection alive across idle intervals without dispatching a client-side event:

```

1 from flama.http.responses import ServerSentEvent,
   ServerSentEventResponse
2
3 async def token_stream():
4     for token in model.generate_tokens(prompt):
5         yield ServerSentEvent(event="token", data=token,
6                               id=str(seq))
7
8 @app.get("/stream")
9 async def sse_endpoint():
10     return ServerSentEventResponse(token_stream())

```

SSE is the primary streaming format for the LLM serving subsystem’s OpenAI, Anthropic, and native dialects. Client libraries can resume interrupted streams via the `Last-Event-ID` header, and the native dialect leverages this for transparent reconnection with sequence-based replay.

Newline-Delimited JSON (NDJSON).. The `NDJSONResponse` (NDJSON Working Group, 2014) emits one compact JSON object per line with the Content-Type set to `application/x-ndjson`. Each yielded item is encoded as a self-contained JSON line followed by a newline character:

```

1 from flama.http.responses import NDJSONResponse
2
3 async def generate_events():
4     async for event in engine.stream():
5         yield event.to_dict()
6
7 @app.get("/events")
8 async def ndjson_endpoint():
9     return NDJSONResponse(generate_events())

```

NDJSON is the streaming format used by the Ollama dialect. Unlike SSE, NDJSON does not support named event types or client reconnection semantics, but its simplicity makes it well-suited for log-style streaming where each line is a self-contained record.

Combined with WebSocket endpoints, these three streaming response types give FLAMA full support for real-time communication patterns: WebSocket for bidirectional messaging, SSE for structured event streams with reconnection, NDJSON for line-oriented progressive delivery, and **StreamingResponse** for arbitrary binary or text streaming.

4.6 OpenAPI schema generation

Every registered route contributes to an OpenAPI 3.2.0 specification ([OpenAPI Initiative, 2025](#)) that is built at startup by the **SchemaModule**. The generation process is automatic and requires no additional annotations beyond the handler's type signature:

- Path parameters become OpenAPI path parameter objects, with types inferred from the route pattern's type converters.
- Query parameters are detected from handler parameters that are not path parameters and not schema objects.
- Request bodies are generated from handler parameters whose type is a registered schema class.
- Response schemas are generated from the handler's return type annotation.
- Schema classes are converted to reusable **components/schemas** definitions via the schema adapter's JSON Schema emission.

Handlers can supply additional metadata (tags, summaries, descriptions, and response codes) as YAML in their docstrings. The docstring is split on `--` and its last segment is parsed with `yaml.safe_load()`, then merged into the generated OpenAPI operation object. The separator is therefore only needed when the docstring opens with human-readable prose that should not be parsed as YAML; a docstring that is entirely YAML, as below, needs no separator:

```
1 @app.get("/users/{user_id:int}")
2 async def get_user(user_id: int) -> UserResponse:
3     """
4     tags:
5         - Users
6     summary:
7         Retrieve a user by ID
8     responses:
9         200:
10             description: The requested user
11         404:
12             description: User not found
13     """
14     ...
```

The generated specification is served as JSON at `/schema/` and rendered as an interactive Swagger UI at `/docs/`.

5 Schema and data validation

Validation is the business of checking that incoming data has the structure and types it claims before any of it reaches application logic. For a web API that means path parameters, query strings, request

bodies, and, if the developer wants it, response bodies too. When something fails the check the client deserves an error it can act on, which in FLAMA's case is a 400 carrying a per-field `detail` object rather than a bare status line. Skip the check and malformed input travels inward instead, surfacing later as a cryptic runtime error, as corrupted data, or as a security problem.

Four things make up the machinery: the adapter that keeps the framework independent of any one schema library, the internal schema representation everything else is written against, the validation pipeline itself, and response serialization.

5.1 The adapter pattern

Python's ecosystem offers several mature libraries for data validation and serialization. Each library has a distinct API, a different approach to field declaration, and its own trade-offs between performance, flexibility, and ease of use. Coupling a web framework to a single validation library forces users to adopt that library's conventions for their entire codebase, even if another library is better suited to their domain or their team's existing investment.

FLAMA avoids this coupling by interposing an adapter layer between the framework's validation pipeline and the schema library. The adapter normalizes three operations:

1. **Field introspection:** given a schema class, extract the list of fields, their types, their nullability, their default values, and whether they are required.
2. **Validation:** given a dictionary of raw data and a schema class, validate the data and return either a validated object or a list of per-field errors.
3. **JSON Schema emission:** given a schema class, produce a JSON Schema object suitable for inclusion in an OpenAPI specification.

Three adapters are provided:

Library	Characteristics
Pydantic (Colvin, 2017)	The most widely adopted validation library. Rust-accelerated core validation. Declarative field definitions with Python type annotations. First-class JSON Schema support.
Marshmallow (Loria, 2013)	A mature library with explicit field declarations, a rich plugin ecosystem, and a serialization API oriented around <code>load()/dump()</code> semantics.
Typesystem (Christie, 2019)	A lightweight library focused on data validation and form rendering. Suitable for small projects where Pydantic or Marshmallow would be over-specified.

The adapter is selected at application construction time and applies globally:

```
1 app = Flama(schema_library="pydantic")
```

Users define schemas using their chosen library's native API (e.g. Pydantic `BaseModel` subclasses, Marshmallow `Schema` subclasses) and reference them in handler signatures. The framework translates between the library's types and its own internal model transparently.

5.2 Internal schema representation

The framework's core validation pipeline operates on an internal schema representation that is independent of any external library. This representation consists of two frozen data classes:

Field represents a single typed field. It records the field's name, Python type, nullability, whether the field is required, its default value (if any), whether it accepts multiple values (i.e. a list), and its JSON Schema equivalent. Fields are the atomic unit of validation: each field is validated independently, and per-field error messages are collected into the validation error response.

Schema is a named collection of **Field** objects. A schema can be constructed from a type annotation (`Schema.from_type(UserUpdate)`) or from an explicit field list (`Schema.build("UserUpdate", fields=[...])`), enabling programmatic schema construction for generated resources.

This internal representation serves as the lingua franca between the adapter layer and the rest of the framework. The validation pipeline, the OpenAPI generator, and the response serializer all consume **Field** and **Schema** objects rather than library-specific types.

5.3 Request validation pipeline

Validation in FLAMA is not a standalone function that the handler calls manually. It is implemented as a sequence of dependency-injection components (Section 6), each responsible for validating one layer of the incoming request. The components execute automatically during the dependency resolution phase (Stage 4 of the request pipeline), before the handler is called.




Figure 3 Request validation pipeline. Each component resolves one aspect of the incoming data and passes validated, typed values to downstream components or the handler.

Four components carry a typical body-validating request through the pipeline, described below. Two siblings not shown in the figure, **PrimitiveParamComponent** and **FileParamComponent**, resolve primitive-typed path and query parameters and **UploadFile** parameters respectively, using the same mechanism.

1. **RequestDataComponent**. This component reads the **Content-Type** header and selects the appropriate codec: `JSONDataCodec` for `application/json`, `URLEncodedCodec` for `application/x-www-form-urlencoded`, and `MultiPartCodec` for `multipart/form-data`. The codec parses the raw request body into `types.RequestData`. If no codec matches the content type, a `NoCodecAvailable` error is raised.
2. **ValidatePathParamsComponent**. This component receives the path parameters extracted by the route table and validates them against the types declared in the route pattern. For example, if the route pattern is `/users/{user_id:int}`, this component verifies that the extracted value is a valid integer.

3. **ValidateQueryParamsComponent.** This component validates query string parameters against the types declared in the handler signature. Parameters that are not path parameters and not schema objects are assumed to be query parameters.
4. **CompositeParamComponent.** This is the component that actually binds a value to a handler parameter annotated `t.Annotated[schemas.Schema, schemas.SchemaMetadata(X)]` (or `list[schemas.Schema]` for a collection). It reads the parsed body produced by `RequestDataComponent` and validates it against `X` through the schema adapter's `validate()` method, raising a `SchemaValidationError` on failure. A separate `ValidateRequestDataComponent` performs the same validation and is available for a handler that wants the raw validated dictionary directly, by declaring a parameter of type `ValidatedRequestData`, without binding it to a specific schema-annotated parameter.

If any validation component encounters an error, the pipeline short-circuits and the framework returns a 400 Bad Request response with a structured JSON body. The `detail` object is keyed by field name; the shape of each field's error entry comes directly from the configured schema library's own error representation. With the Pydantic adapter, for example:

```
{
  "status_code": 400,
  "detail": {
    "age": {
      "type": "int_parsing",
      "loc": ["age"],
      "msg": "Input should be a valid integer, unable to
              parse string as an integer",
      "input": "not-an-int"
    }
  },
  "error": "ValidationError"
}
```

5.4 Response serialization

Return values from handlers are serialized through the `APIResponse` class. When the route declares a response schema, `APIResponse` re-validates the returned value through the schema adapter, which reconstructs the schema object from the value's keys and dumps it back to a plain dictionary; the returned value must therefore already be a dictionary (or a list of dictionaries) keyed by the schema's field names, not an instantiated schema class. The resulting dictionary is then encoded to JSON using the Rust-accelerated JSON encoder.

The response type and its JSON Schema are inferred from the handler's return type annotation and contribute to the OpenAPI specification. This means that the return type annotation has two purposes: it controls the runtime serialization of the response, and it generates the documentation that clients use to understand the response format.

6 Dependency injection

Dependency injection is how a handler's parameters get resolved without the handler having to build or find its own dependencies (Fowler, 2004). In FLAMA it is not an optional convenience layered

over something simpler. It is the mechanism the framework itself uses to hand over request data, validated inputs, database connections, authentication tokens, and model instances, so every request makes at least one pass through the injector whether the developer engages with it or not.

What follows covers the component model, the resolution tree built at startup, the per-request cache, and the context types available without registering anything.

6.1 The component model

A **component** is a class that knows how to produce a value of a given type. It is the fundamental building block of FLAMA's dependency injection system. Every component implements two methods:

1. `can_handle_parameter(parameter)`: inspects a parameter's type annotation and returns `True` if the component can produce a value for that parameter. The default implementation checks whether the parameter's type matches the return type of the component's `resolve()` method.
2. `resolve(**kwargs)`: produces the value. The method's return type annotation indicates the type of the produced value, and the method's parameter annotations indicate the component's own upstream dependencies. These dependencies are resolved recursively using the same injection mechanism.

The following example defines a component that produces an asynchronous database connection:

```
1 from flama import Component
2
3 class DatabaseConnectionComponent(Component):
4     async def resolve(self, app: Flama) -> AsyncConnection:
5         return await app.sqlalchemy.open_connection()
```

The component declares that it requires the `Flama` application instance (a context value, always available) and produces an `AsyncConnection`, which it obtains from the `SQLAlchemyModule` registered on the application and reached, like any module, under its name. When a handler declares a parameter of type `AsyncConnection`, the injector identifies this component as the provider and calls its `resolve()` method with the application instance.

Components can depend on other components, forming a dependency graph of arbitrary depth:

```
1 class UserRepositoryComponent(Component):
2     def resolve(self,
3                 connection: AsyncConnection) -> UserRepository:
4         return UserRepository(connection)
```

Here, `UserRepositoryComponent` depends on `AsyncConnection`, which is itself resolved by `DatabaseConnectionComponent`. The injector resolves the full chain automatically.

6.2 Resolution trees

At startup, the injector examines every registered handler and builds a **resolution tree** for each of its parameters. This tree is a directed acyclic graph (DAG) where each node represents a value that must be produced, and each edge represents a dependency relationship.

Each node in the tree is one of three types:

ContextNode the value is available directly from the ASGI scope or the request context. Context values include the `Request` object, the `Flama` application, the matched `Route`, and the `WebSocket` connection (for `WebSocket` handlers).

ComponentNode the value is produced by a registered component. The node stores a reference to the component instance and has child nodes for the component's own dependencies.

ParameterNode the value is a primitive (string, integer, float, boolean, UUID) extracted from a path segment or query string parameter.




Figure 4 Resolution tree for a handler with four parameters. The injector traverses the tree bottom-up: context values are read from the ASGI scope, component values are produced by calling `resolve()`, and primitive parameters are extracted from the URL path. The tree is compiled at startup and flattened into an ordered sequence of resolution steps for efficient request-time execution. Node fill indicates where a value comes from, as given in the key.

Figure 4 shows the resolution tree for a handler with four parameters: a path parameter, a schema validation result, a database connection (component), and an authentication token (component). The injector compiles this tree at startup and flattens it into an ordered sequence of **resolution steps** that can be executed without graph traversal at request time. Circular dependencies are detected during compilation and raise a descriptive error.

6.3 Caching

Components may declare `cacheable = True` (the default), in which case their resolved value is stored in a per-request LRU cache keyed by the component's identity string. The cache ensures that a database connection, an authentication token, or any other expensive-to-produce value is resolved at most once per request, even if multiple handler parameters or nested components depend on it.

The cache is scoped to a single request and is discarded after the response has been sent. This prevents stale values from leaking across requests and ensures that each request sees a fresh set of dependencies.

6.4 Context types

The following types are available as context values without requiring a component. Any handler or component can request them by declaring parameters with these type annotations:

Type	Value
Scope	The ASGI scope dictionary, containing the request method, path, headers, query string, and client address.
Receive	The ASGI receive channel, from which request body chunks are read.
Send	The ASGI send channel, through which response headers and body are emitted.
Request	The HTTP request object, with properties for method, path, query parameters, headers, cookies, and methods for reading the body.
Response	The HTTP response object. Available in middleware and post-processing hooks.
WebSocket	The WebSocket connection object. Available only in WebSocket handlers.
App	The application instance (Flama). Provides access to the application's state, configuration, modules, and registered routes.
Route	The matched route object. Contains the route pattern, allowed methods, and handler reference.
Exception	The current exception, if the handler is being invoked in an error-handling context.

7 Resources and domain-driven design

A SQLAlchemy table and a schema class are enough, between them, to generate nine REST endpoints with correct status codes, pagination, and error handling, backed by the domain-driven design patterns that keep the transactions honest. This is the part of FLAMA that saves the most typing, and it is also the part that most needs an escape hatch, since a generated CRUD API is only useful for as long as it does what the domain wants. Overriding a single operation is therefore possible without giving up the other eight.

7.1 The resource abstraction

A **resource** in FLAMA is a class that declares the data model, the validation schema, and the set of operations that an API exposes over a domain entity. The framework provides two resource base classes:

Resource a generic resource for virtual entities that do not correspond to a database table: health checks, calculators, proxies to external APIs, or aggregation endpoints. The developer defines the operations explicitly.

CRUDResource a database-backed resource that derives its operations from a SQLAlchemy ([Bayer, 2006](#)) table definition. The framework's metaclass inspects the table columns, generates input and output schemas, creates a repository class, and produces handler methods for all standard CRUD operations.

7.2 CRUD resource declaration

A **CRUDResource** is declared by providing three pieces of information: the SQLAlchemy table, the schema class for output (and optionally separate schemas for input), and a human-readable name:

```

1 import sqlalchemy
2 from flama.resources.crud import CRUDResource
3
4 metadata = sqlalchemy.MetaData()
5
6 users = sqlalchemy.Table(
7     "users",
8     metadata,
```

```

9     sqlalchemy.Column("id", sqlalchemy.Integer,
10                        primary_key=True,
11                        autoincrement=True),
12     sqlalchemy.Column("name", sqlalchemy.String(100),
13                        nullable=False),
14     sqlalchemy.Column("email", sqlalchemy.String(200),
15                        nullable=False, unique=True),
16     sqlalchemy.Column("created_at", sqlalchemy.DateTime,
17                        server_default=sqlalchemy.func.now()),
18 )
19
20 class UserResource(CRUDResource):
21     name = "user"
22     verbose_name = "User"
23     model = users
24     schema = UserSchema
25     input_schema = UserInputSchema
26     output_schema = UserOutputSchema

```

The class declaration only defines the resource; it still has to be attached to the application, which mounts its generated routes under the given path prefix:

```

1 app.resources.add_resource("/users/", UserResource)

```

`CRUDResource` operates on an async SQLAlchemy connection, so the application also needs the `SQLAlchemyModule` (Section 7.3) registered to actually have a database engine to connect to; without it, the generated handlers have nothing to run their queries against.

From this declaration, the metaclass generates the following endpoints:

Method	Path	Status	Behaviour
POST	/users/	201	Validate input, insert row, return created resource.
GET	/users/{resource_id}/	200	Retrieve single resource by primary key, or 404.
PUT	/users/{resource_id}/	200	Full replacement of a single resource.
PATCH	/users/{resource_id}/	200	Partial update of a single resource.
DELETE	/users/{resource_id}/	204	Delete single resource, or 404.
GET	/users/	200	List resources with pagination.
PUT	/users/	200	Bulk replace all resources.
PATCH	/users/	200	Bulk partial replace of all resources.
DELETE	/users/	204	Bulk delete all resources (drop), reporting the number deleted.

The single-resource routes are keyed on `resource_id` rather than the column's own name, which is also the keyword that `app.resolve_url()` expects when building a URL for one of them (e.g. `app.resolve_url("user:retrieve", resource_id=1)`).

Each operation includes error handling: integrity errors (e.g. a duplicate email violating a unique constraint) are caught as `IntegrityError` and produce 400 Bad Request responses, and missing resources produce 404 Not Found responses. Custom methods can be added using the `@ResourceRoute` `.method()` decorator, and any generated method can be overridden by defining it in the subclass.

7.3 Domain-driven design patterns

The generated CRUD endpoints do not interact with the database directly. Instead, they delegate data access and transactional management to two domain-driven design patterns (Evans, 2003; Fowler, 2002): the Repository and the Unit of Work (called Worker in FLAMA).

7.3.1 The Repository pattern

A **repository** encapsulates all interaction with a data source behind a collection-like interface. The calling code does not know whether the data comes from a relational database, an in-memory cache, or a remote HTTP service. This separation has three benefits: it makes the data access logic testable in isolation, it allows the data source to be changed without modifying the business logic, and it provides a clear boundary for caching and query optimization.

FLAMA provides two repository families:

SQLAlchemyTableRepository operates on a SQLAlchemy table using async connections. Every method other than `create` locates rows through SQLAlchemy clauses and exact-match filters passed as `*clauses/**filters`, the same vocabulary as a SQLAlchemy `select()`, rather than a bare primary key. It provides six methods: `create(*data)` is variadic, inserting one row per positional dict and always returning a list of the created rows; `retrieve(*clauses, **filters)` fetches a single row, raising if none or more than one match; `update(data, *clauses, **filters)` updates the matching rows with `data` and returns them; `delete(*clauses, **filters)` removes the matching rows; `list(*clauses, order_by=None, order_direction="asc", **filters)` returns an async iterator over the matching rows (compatible with the pagination system); `drop(*clauses, **filters)` removes the matching rows and returns how many were dropped.

HTTPResourceRepository proxies CRUD operations over HTTP to a remote service, so one FLAMA application can consume resources exposed by another while its calling code still speaks in repository terms. A subclass declares the remote resource's path (`_resource = "/user"`) and is constructed with a `Client` pointing at the service. Because it addresses a REST resource rather than a table, it is keyed on identity rather than on query clauses: `create(data)`, `retrieve(id)`, `update(id, data)`, `partial_update(id, data)`, and `delete(id)` act on a single record, while `list()`, `replace(data)`, `partial_replace(data)`, and `drop()` act on the collection. `list()` follows the remote endpoint's pagination transparently, yielding records across pages as an async iterable.

The two families therefore share a vocabulary rather than a literal signature: the SQLAlchemy repository selects rows with clauses and filters, while the HTTP repository addresses resources by identifier, because that is what each underlying data source exposes.

7.3.2 The Unit of Work pattern (Worker)

A **Worker** is the FLAMA implementation of the Unit of Work pattern. The transactional boundary logic (connection acquisition, begin, commit, and rollback) is defined by the abstract base class `AbstractWorker`, which implements the async context manager protocol (`__aenter__`/`__aexit__`). The concrete `Worker` subclass provides default (no-op) implementations of `set_up()`, `tear_down()`, `commit()`, and `rollback()`, suitable for applications that override these methods with database-specific behaviour. All repository operations within a Worker's context execute atomically: either all operations commit, or all are rolled back.

In practice, a `Worker` for a SQLAlchemy-backed application subclasses `SQLAlchemyWorker` rather than the bare `Worker`, which supplies the `set_up()/tear_down()/commit()/rollback()` implementations that talk to a SQLAlchemy connection:

```
1 from flama.ddd.workers.sqlalchemy import SQLAlchemyWorker
2
3 class AppWorker(SQLAlchemyWorker):
4     users: UserRepository
5     posts: PostRepository
```

The Worker's type annotations declare which repositories it manages; repository instances are created lazily from these annotations when the Worker is used as an async context manager. A `WorkerComponent` makes a specific Worker instance available for dependency injection, so a handler that declares a parameter of that Worker's type receives a fully configured instance; the component still has to be registered explicitly, alongside the `SQLAlchemyModule` that gives it a database engine to connect to:

```
1 from flama import Flama
2 from flama.ddd import WorkerComponent
3 from flama.sqlalchemy import SQLAlchemyModule
4
5 app = Flama(
6     modules=[SQLAlchemyModule(DATABASE_URL)],
7     components=[WorkerComponent(worker=AppWorker())],
8 )
9
10 @app.post("/users/")
11 async def create_user(worker: AppWorker) -> None:
12     async with worker:
13         # create() is variadic and always returns a list, one
14         # created row per positional dict passed in.
15         [user] = await worker.users.create(
16             {"name": "Ada", "email": "ada@example.com"}
17         )
18         await worker.posts.create(
19             {"author_id": user["id"],
20              "title": "First post",
21              "body": "Hello, world."}
22         )
23         # Commits on successful exit.
24         # Rolls back on exception.
```

The Worker manages four lifecycle operations: connection acquisition, transaction begin, commit (on successful exit of the `async with` block), and rollback (on exception).

8 Authentication and authorization

FLAMA ships its own JWT ([Jones et al., 2015b](#)) implementation rather than reaching for a third-party library, and the reason is integration rather than distrust of the alternatives. A token that the framework decodes itself can be handed to a handler as a typed parameter by the same injector that supplies everything else, checked by middleware that already knows how routes are tagged, and




Figure 5 Domain-driven design integration. The handler receives a Worker via dependency injection. The Worker manages repositories within a transactional boundary. All operations within the `async with` block either commit together or roll back together.

described in the generated OpenAPI document without a separate registration step. Three pieces follow: the JWT implementation, the components that make a decoded token available to handlers, and the middleware that turns its claims into route-level access control.

8.1 JWT implementation

The framework provides a complete JWT implementation covering the three aspects of token-based authentication:

Token encoding and decoding. The `JWT` class encodes a payload dictionary into a signed token string and decodes a token string back into a payload, verifying the signature in the process. Standard claims (`iss`, `sub`, `aud`, `exp`, `iat`, `nbf`) are validated automatically on decoding. Expired tokens and tokens with future `nbf` claims are rejected.

Signing algorithms. HMAC-SHA256 (HS256) is the default signing algorithm. HS384 and HS512 are also implemented, for applications that require a longer MAC. All three are symmetric: the signing key is shared between the token issuer and the token consumer. Asymmetric algorithms (RSA, ECDSA) are not implemented, so issuer and consumer must currently share the same secret key.

JSON Web Signature (JWS). The compact serialization format is implemented following the JSON Web Signature specification (Jones et al., 2015a). A token consists of three Base64-URL-encoded segments separated by dots: the header (algorithm and token type), the payload (claims), and the signature.

8.2 Token components

Authentication integrates with the dependency injection system through two components, each of which resolves a typed token object from the incoming request:

AccessTokenComponent extracts the JWT from an `access_token` header or an `access_token` cookie (header checked first), by default expecting the header value in the form `Bearer <token>`.

It decodes the token against the secret it was constructed with and returns an `AccessToken` object containing the decoded header and payload.

RefreshTokenComponent performs the same extraction and decoding for refresh tokens, reading a `refresh_token` header or cookie instead. Refresh tokens are used in token rotation flows where the client exchanges an expired access token and a valid refresh token for a new pair of tokens.

Neither component is registered by default: both take the signing secret as a required constructor argument, so the application declares the ones it needs explicitly, alongside the middleware that enforces permissions on top of them (Section 8.3):

```
1 from flama import Flama
2 from flama.authentication import AccessTokenComponent,
   AuthenticationMiddleware
3
4 app = Flama(
5     components=[AccessTokenComponent(secret=SECRET_KEY)],
6     middleware=[AuthenticationMiddleware()],
7 )
```

Any handler that needs to know the identity of the authenticated user simply declares a parameter of type `AccessToken`:

```
1 import typing as t
2 from flama import schemas
3 from flama.authentication import AccessToken
4
5 ProfileResponse = t.Annotated[
6     schemas.Schema, schemas.SchemaMetadata(Profile)
7 ]
8
9 @app.get("/profile")
10 async def get_profile(token: AccessToken) -> ProfileResponse:
11     user_id = token.payload.sub
12     return await load_profile(user_id)
```

The framework resolves the token transparently. If the request does not contain a valid token, the component raises an exception that the middleware translates into a 401 Unauthorized response. Standard claims (`iss/sub/aud/exp/nbf/iat/jti`) are attributes of `token.payload`; anything else, such as application-specific user data, is namespaced under `token.payload.data`, a plain dictionary, since it falls outside the JWT standard and is not validated on decoding.

8.3 Permission-based middleware

The `AuthenticationMiddleware` enforces route-level access control based on JWT permission claims. Routes are tagged with required permissions at registration time:

```
1 @app.route("/admin/users", methods=["DELETE"],
2           tags={"permissions": ["admin", "delete"]})
3 async def delete_all_users() -> None:
4     ...
```

When a request arrives at a route whose tags declare required permissions, the middleware performs the following steps:

1. Resolve an `AccessToken` for the request, the same way a handler parameter of that type would be resolved.
2. Read the user's permissions from `token.payload.data["permissions"]` (a list of strings), unioned with every permission listed under each role in `token.payload.data["roles"]` (a mapping of role name to its list of permissions), so a client can be granted access directly or through a role.
3. Verify that the user's permissions are a superset of the route's required permissions.
4. If the check passes, forward the request to the handler. If the token is missing or invalid, return 401 Unauthorized. If the token is valid but the permissions are insufficient, return 403 Forbidden.

The tag key the middleware reads (`permissions` by default) and a list of URL regex patterns to exempt from authentication entirely are both configurable on `AuthenticationMiddleware`'s constructor.

This design separates authentication (verifying identity) from authorization (verifying permissions): the token components handle authentication, and the middleware handles authorization. The two concerns can be used independently: a handler can request an `AccessToken` without the middleware being active, and the middleware can enforce permissions without the handler needing to inspect the token.

9 Pagination

A list endpoint over a table of any size eventually has to answer the question of what to leave out. Serialising a million rows into one JSON response wastes bandwidth, delays the first byte, and may simply exhaust the client, so what is wanted instead is a way for the caller to ask for a slice and a standard envelope telling it which slice arrived. FLAMA offers two strategies for this. Both work by wrapping the list handler, adding their query parameters to its signature without the handler's code being aware of them, and folding the returned collection into the envelope on the way out.

9.1 Pagination strategies

Page-number pagination. The client specifies a `page` (1-indexed) and a `page_size`. This strategy is natural for tabular UIs where the user navigates between numbered pages.

Limit-offset pagination. The client specifies a `limit` (maximum number of items) and an `offset` (number of items to skip). This strategy is suited to infinite-scroll interfaces and cursor-based navigation.

Both strategies support an optional `count` query parameter (boolean, default `False`) that controls whether the total number of items is reported. It is off by default because counting is the expensive part of paginating a large table; when omitted, the envelope still carries the pagination parameters but reports `null` for the count.

9.2 Usage

Pagination is applied declaratively at the route level:

```

1 @app.route("/items/", pagination="page_number")
2 async def list_items(**kwargs) -> list[Item]:
3     return await repository.list()

```

The `**kwargs` is not decoration. The paginator rewrites the handler's signature to add its own query parameters, and it refuses to wrap a handler that has nowhere to put them, raising `TypeError` at registration time if the parameter is absent.

The paginator intercepts the handler's return value, slices the collection according to the pagination parameters, and wraps the result in a response envelope:

A request to `/items/?page=2&page_size=20&count=true` then produces:

```

{
  "meta": {
    "page": 2,
    "page_size": 20,
    "count": 148
  },
  "data": [
    {"id": 21, "name": "Widget A"},
    {"id": 22, "name": "Widget B"}
  ]
}

```

The `meta` object provides the pagination parameters used for the current request and the total count (null unless `count=true` was requested). For limit-offset pagination, the `meta` object contains `limit`, `offset`, and `count` fields instead. The default page size is 10.

The paginator is implemented as a handler wrapper. It replaces the handler's `__signature__` with one in which `**kwargs` has been substituted by the strategy's own parameters, which is why the framework can resolve and document them without the handler ever naming them. It then intercepts the raw return value and delegates slicing and counting to the appropriate strategy class. Because the substitution happens on the signature the rest of the framework reads, the pagination parameters also appear in the auto-generated OpenAPI specification.

10 Background tasks

Sending a confirmation email, generating a report, writing an audit log, kicking off a retraining run: none of these should hold up the HTTP response, and all of them have to happen somewhere. A synchronous framework leaves two options, neither good, which are to make the caller wait or to stand up an external queue such as Celery for what may be a single line of work. FLAMA runs them itself, after the response body has gone out, with no additional infrastructure involved.

10.1 Concurrency models

Two execution strategies are available, corresponding to the two dominant forms of concurrency in Python:

BackgroundThreadTask executes the callable in a thread via `asyncio.to_thread()`. This model is appropriate for I/O-bound work: sending HTTP requests, writing to a message queue, or

performing database operations.

BackgroundProcessTask executes the callable in a separate process via `multiprocessing.Process`. This model is appropriate for CPU-bound work: model retraining, image resizing, PDF generation, or any computation that would block the event loop if run in a thread.

10.2 Usage

Tasks are created and attached to responses:

```
1 import typing as t
2 from flama import BackgroundThreadTask, schemas
3 from flama.http import APIResponse
4
5 OrderRequest = t.Annotated[
6     schemas.Schema, schemas.SchemaMetadata(OrderInput)
7 ]
8
9 @app.post("/orders/")
10 async def create_order(data: OrderRequest):
11     order = await save_order(data)
12     task = BackgroundThreadTask(
13         send_confirmation_email,
14         recipient=order["email"],
15         order_id=order["id"],
16     )
17     return APIResponse(order, status_code=201,
18                        background=task)
```

The response is sent immediately. The email is sent in a background thread after the response body has been fully transmitted.

For handlers that need to schedule multiple tasks, the `BackgroundTasks` class aggregates them into a single container:

```
1 from flama import BackgroundTasks
2
3 tasks = BackgroundTasks()
4 tasks.add_task("thread", send_email, order["email"])
5 tasks.add_task("process", generate_invoice, order["id"])
6
7 return APIResponse(order, background=tasks)
```

Tasks in a `BackgroundTasks` container are executed sequentially in the order they were added.

11 Middleware

Logging, error handling, authentication, compression, CORS: concerns that belong to every request and to no handler in particular. Middleware is where they live. In FLAMA it is an ordered stack of ASGI callables each wrapping the next, on the onion model that ASGI and WSGI frameworks have converged on, where a request enters at the outermost layer, works inward through each wrapper to the handler, and the response comes back out through the same layers in reverse. Any layer may

inspect or modify what passes through it in either direction, or answer immediately and let nothing further in.

11.1 Middleware stack

Construction is bottom-up, so the middleware registered last ends up outermost and sees each request first. Worth keeping in mind when order matters: authentication that must run before compression has to be registered after it.

Depth costs less than one might expect. Measured under Callgrind, a request through an application with no middleware at all costs 3.52 M estimated cycles, with five 3.53 M, and with ten 3.54 M, so each additional layer adds on the order of 0.1%. The stack is cheap enough that the decision about what to put in it can be made on architectural grounds rather than on a per-request budget.

11.2 Built-in middleware

FLAMA ships with the following middleware:

Middleware	Purpose
BaseHTTPMiddleware	A convenience class for writing HTTP middleware using a simple <code>dispatch(request, call_next)</code> interface instead of raw ASGI callables.
ExceptionMiddleware	Maps Python exceptions to HTTP responses. Supports custom handlers per exception type. In debug mode, renders interactive HTML error pages with full tracebacks, source code context, and local variable values.
ServerErrorMiddleware	Catches any exception that escapes the <code>ExceptionMiddleware</code> . Produces a generic 500 Internal Server Error response and logs the traceback.
CORSMiddleware	Adds Cross-Origin Resource Sharing headers. Configurable allowed origins, methods, and headers. Handles preflight <code>OPTIONS</code> requests automatically.
CompressionMiddleware	Negotiates a compression codec from the client's <code>Accept-Encoding</code> header and compresses the response body. Brotli and gzip are tried by default, in that order; a configurable minimum response size gates when compression is attempted.
HTTPSRedirectMiddleware	Redirects HTTP requests to HTTPS with a 301 permanent redirect.
TrustedHostMiddleware	Validates the <code>Host</code> header against a whitelist of allowed hosts. Returns 400 for requests with untrusted hosts.
SessionMiddleware	Provides cookie-based HTTP sessions. Session data is serialized to JSON and signed as a JWS token with HMAC-SHA256, with expiry enforced from the token's issued-at claim.
CorrelationIdMiddleware	Assigns a correlation ID to every request, taken from an incoming header or generated as a UUID4, and echoes it back on the response for downstream log correlation.
AuthenticationMiddleware	Enforces JWT permission-based access control (Section 8). Defined in the <code>flama.authentication</code> module rather than the core middleware module.

11.3 Custom middleware

Custom middleware subclasses `Middleware` and implements the ASGI `__call__`. The downstream application is not passed to the constructor; `MiddlewareStack` injects it as `self.app` when it assembles the chain, which leaves `__init__` free to take the middleware's own configuration:

```
1 import time
2 import logging
3
```

```

4 from flama.middleware import Middleware
5
6 logger = logging.getLogger(__name__)
7
8 class TimingMiddleware(Middleware):
9     async def __call__(self, scope, receive, send):
10         if scope["type"] != "http":
11             await self.app(scope, receive, send)
12             return
13
14         start = time.monotonic()
15         await self.app(scope, receive, send)
16         duration = time.monotonic() - start
17         path = scope.get("path", "/")
18         logger.info("%s completed in %.3fs",
19                     path, duration)

```

Middleware is registered at application construction time:

```

1 from flama import Flama
2 from flama.middleware import CORSMiddleware
3
4 app = Flama(
5     middleware=[
6         TimingMiddleware(),
7         CORSMiddleware(allow_origins=["*"]),
8     ],
9 )

```

What the stack receives are middleware *instances*, already configured, and it wires each one to its downstream neighbour during startup. Earlier versions took the class and its arguments separately, wrapped in a `Middleware(...)` call; that form was removed in 2.0 and now raises `TypeError`.

Part III

Machine Learning and Generative AI

12 Machine-learning model serving

The central premise of FLAMA is that the gap between training a machine-learning model and deploying it behind a production API should be as small as the gap between writing a function and exposing it as an HTTP endpoint. In practice, this gap is often the dominant source of friction in ML projects: the model is trained in a notebook or a training script, exported as a pickle file or a checkpoint directory, and then handed to a backend engineer who must write serialization code, input validation, error handling, health checks, and deployment scripts from scratch. Closing that gap takes four things, and they make up the rest of this part: a portable binary format for the serialized model, an integration with dependency injection so the loaded model behaves like any other dependency, a resource abstraction that turns one class declaration into a working endpoint, and serializers for the four frameworks that produce most models in practice.

The order is bottom-up. The binary format that holds a model on disk comes first, then the framework-specific serializers that read and write it, then the mechanics of getting a serialized model into a running application, and last the three levels of integration on offer, which run from a single predict endpoint to a full resource with inspection and streaming.

12.1 Design goals

Four requirements shaped the design of the model serving subsystem:

1. **Framework agnosticism.** The deployment system must support models trained with scikit-learn, TensorFlow/Keras, PyTorch, and HuggingFace Transformers. Adding a new framework should require only a serializer and a model wrapper, not changes to the serving infrastructure.
2. **Self-describing artifacts.** A deployed model must carry its own metadata: the framework and version that produced it, the model's class name, its hyperparameters, its training metrics, and any auxiliary artifacts (tokenizer files, vocabulary files, label maps). This metadata must be accessible without loading the model weights, enabling fast inspection and cataloguing.
3. **Efficient transport.** Model files are frequently large, hundreds of megabytes for a convolutional network and gigabytes for a language model, so the format has to compress without making the load path expensive. The asymmetry is deliberate and shows up in measurement: for a scikit-learn artifact under protocol 2, a dump costs around 43.7 M estimated cycles against 6.4 M for the corresponding load. Packaging happens once and loading happens on every cold start, so that is the right way round.
4. **Zero-configuration serving.** Given a model file, it should be possible to expose it as a REST endpoint with a single command and no application code. The framework should generate input/output schemas, OpenAPI documentation, and error handling automatically.

13 The FLM binary format

The `.flm` format (for **F**lama **L**earned **M**odel) packages a serialized model, its metadata, and whatever auxiliary artifacts it needs into one self-describing file. Two properties drove the design. It is versioned, because the storage requirements of a scikit-learn pickle and a multi-file LLM checkpoint have nothing in common and a single layout would have had to be stretched to cover both. And the metadata sits at a known offset ahead of the weights, so reading it costs a header parse rather than a full decompression, which is what lets the framework decide how to load a multi-gigabyte artifact before committing to loading it.

13.1 File structure

An FLM file consists of a 16-byte outer header followed by a body whose layout is determined by the protocol version declared in the header. Two protocol versions are defined: version 1 for traditional ML models serialized as opaque binary blobs, and version 2 for models that may be stored as directory bundles (such as LLM checkpoints) and that benefit from per-section compression control.




Figure 6 Structure of the FLM binary format (protocol version 1). The outer header identifies the format version and compression algorithm. The body contains three sections: compressed JSON metadata, compressed model weights, and a sequence of compressed named artifact entries. Each section can be decompressed independently.

13.1.1 Outer header

The outer header uses network byte order (**big-endian**) and consists of three fields packed according to the `struct` format `!I I Q`:

Offset	Size	Field	Description
0	4 B	Protocol version	Unsigned 32-bit integer, 1 or 2. 2 is the default written by <code>dump()</code> .
4	4 B	Compression format	Enum value identifying the compression algorithm.
8	8 B	Body size	Unsigned 64-bit integer. The total size of the body in bytes.

13.1.2 Body header (protocol version 1)

The body begins with a 28-byte header packed as `!Q Q I Q`:

Offset	Size	Field	Description
0	8 B	Meta size	Size of the compressed metadata section.
8	8 B	Model size	Size of the compressed model weights section.
16	4 B	Artifacts count	Number of artifact entries.
20	8 B	Artifacts size	Total size of all compressed artifact entries.

13.1.3 Metadata section

The metadata section is a JSON document compressed with the algorithm specified in the outer header. It is stored at a fixed offset immediately after the body header, which means it can be decompressed and inspected without reading the model weights.

The metadata is represented internally by the `ModelArtifact` data structure, which comprises four nested frozen dataclasses:

```

1  @dataclass(frozen=True)
2  class FrameworkInfo:
3      lib: str          # "sklearn", "tensorflow",
4                        # "torch", "transformers"
5      version: str      # e.g. "1.7.2", "2.20.0"
6
7  @dataclass(frozen=True)
8  class ModelInfo:
9      obj: str          # class name
10     info: dict | None  # JSON Schema (framework-specific)
11     params: dict | None # hyperparameters
12     metrics: dict | None # training metrics
13
14  @dataclass(frozen=True)
15  class Metadata:
16     id: str | UUID      # unique model identifier
17     timestamp: datetime # serialization timestamp
18     framework: FrameworkInfo
19     model: ModelInfo
20     extra: dict | None  # user-defined metadata
21
22  @dataclass(frozen=True)
23  class ModelArtifact:
24     meta: Metadata
25     model: Any          # the deserialized model object
26     artifacts: dict[str, Path] | None

```

13.1.4 Model weights section

The model weights section contains the serialized model, compressed with the same algorithm as the metadata. The serialization format is framework-specific and is described in [Section 14](#).

13.1.5 Artifacts section

The artifacts section is a sequence of named entries. Each entry consists of a 12-byte header (!I Q: 4 bytes for the name length, 8 bytes for the compressed content length), followed by the UTF-8 name string and the compressed content. Artifacts are used for auxiliary files that the model needs

at inference time but that are not part of the weight tensors: tokenizer vocabularies, label maps, preprocessing configurations, and custom post-processing scripts.

On deserialization, artifacts are extracted into a temporary directory that is cleaned up automatically when the `ModelArtifact` object is garbage-collected (via `weakref.finalize`).

13.1.6 Protocol version 2

Protocol version 2 restructures the body layout to accommodate two advances that emerged from the LLM serving subsystem: the need to store model weights as directory bundles (tarballs of Hugging Face checkpoint trees rather than opaque byte blobs) and the desire for per-section compression control so that metadata can remain uncompressed for fast random-access inspection while large model payloads are compressed independently.

The v2 body header is 24 bytes, packed as `!Q Q Q`:

Offset	Size	Field	Description
0	8 B	Meta size	Size of the metadata section (including its discriminator byte).
8	8 B	Artifacts size	Total size of the artifacts section (including its discriminator byte).
16	8 B	Model size	Size of the model section (including its discriminator and kind bytes).

Each section begins with a one-byte *compression discriminator* that declares how that section is compressed:

Byte	Name	Meaning
0x00	inherit	Use the file-level compression declared in the outer header.
0x01	bz2	Override with bz2.
0x02	lzma	Override with lzma.
0x03	zlib	Override with zlib.
0x04	zstd	Override with zstd.
0xFF	none	Passthrough (no compression).

The model section carries an additional one-byte *kind discriminator* immediately after its compression byte:

Byte	Kind	Payload format
0x00	binary	Opaque serialized bytes (pickle, Keras, torch.export). Used for traditional ML models.
0x01	bundle	A tar stream of the model directory. Used for Transformers checkpoints and LLM artifacts, which consist of multiple files (weight shards, tokenizer, configuration).

The combination of per-section compression and model kinds means that an LLM checkpoint can be stored as an uncompressed tar bundle (fast extraction at load time, since model files are already stored in efficient formats like safetensors) while the metadata remains independently accessible and the artifacts section uses a different compression level. The `inherit` discriminator allows sections that do not need special treatment to fall back to the file-level default, maintaining backward compatibility with the compression semantics of protocol version 1.

13.1.7 Model capabilities

Protocol version 2 also introduces a `capabilities` field in the metadata section that declares what a packaged model can ingest and produce. Capabilities are represented by the `ModelCapabilities` hierarchy:

```

1 @dataclass(frozen=True)
2 class LLMModelCapabilities(ModelCapabilities):
3     kind: ClassVar[ModelFamily] = "llm"
4
5     text: bool = True
6     image: bool = False
7     audio: bool = False
8     video: bool = False
9     tools: bool = False
10    reasoning: bool = False

```

Capabilities are detected at serialization time by each framework-specific serializer and persisted in the manifest. Consumers—backend dispatch, input validation, serving-layer capability advertisement—read this single source of truth rather than re-probing the model at load time.

13.1.8 Artifact families

The metadata's `FrameworkInfo` dataclass carries a family discriminator ("ml" or "llm") that records the artifact's intent at serve time. The family is chosen by the producer at dump time (e.g. via `flama get -family llm`) and is never inferred from the library at load time. LLM artifacts always record "transformers" as their library because the on-disk format is a Hugging Face checkpoint tarball; the runtime that actually serves them (vLLM or MLX) is selected at load time by an import probe and is not persisted in the manifest.

13.2 Compression

The FLM format supports four compression algorithms:

Enum value	Algorithm	Notes
1	bz2	Moderate compression ratio, slow decompression.
2	lzma	High compression ratio, very slow decompression.
3	zlib	Balanced compression and speed. Available in all Python installations.
4	zstd (default)	Best trade-off: near-lzma compression ratios at near-zlib decompression speeds. Uses <code>python-zstd</code> or the native <code>compression.zstd</code> module available in Python 3.14+.

The default is zstd ([Collet and Kucherawy, 2021](#)), which provides compression ratios comparable to lzma while decompressing at speeds comparable to zlib. For neural network weights (which consist largely of floating-point tensors with limited redundancy), zstd achieves typical compression ratios of 1.5–3× on serialized model files.

13.3 Serialization and deserialization API

The top-level API consists of two functions:

```

1 from flama.serialize import dump, load
2
3 # Serialize a trained model to disk
4 dump(
5     model,
6     path="classifier.flm",
7     family="ml",

```

```

8     compression="zstd",
9     model_id="classifier-v2",
10    params={"n_estimators": 100,
11            "max_depth": 8},
12    metrics={"accuracy": 0.947,
13            "f1": 0.932},
14    extra={"dataset": "prod-2024-q3",
15          "author": "team-ml"},
16 )
17
18 # Load a serialized model from disk
19 artifact = load(path="classifier.flm")
20
21 print(artifact.meta.framework.lib)      # "sklearn"
22 print(artifact.meta.model.obj)          # "RandomForestClassifier"
23 print(artifact.meta.model.metrics)      # {"accuracy": 0.947, ...}
24 prediction = artifact.model.predict([[5.1, 3.5, 1.4, 0.2]])

```

Both functions accept either a path (string or Path) or a binary file object, making them usable with local files, S3 objects, HTTP responses, or any other binary stream. `dump()` additionally requires a `family` keyword ("ml" or "llm"); it is never inferred from the model object, since the same on-disk `transformers` library backs both a predictive pipeline and an LLM checkpoint.

On loading, the framework version stored in the metadata is compared with the installed version. If they differ, a `FrameworkVersionWarning` is emitted to alert the user that the model was trained with a different version of the framework and that predictions may not be reproducible.

14 Framework-specific serializers

Each supported machine learning framework has a dedicated serializer that implements three operations: `dump()` (model to bytes), `load()` (bytes to model), and `info()` (model to JSON Schema describing the model's architecture or parameters). All serializers inherit from a common abstract base class:

```

1 class BaseModelSerializer(ABC):
2     lib: ClassVar[str]
3
4     @abstractmethod
5     def dump(self, obj, /, **kwargs) -> bytes:
6         """Serialize a model to bytes."""
7         ...
8
9     @abstractmethod
10    def load(self, model: bytes, /, **kwargs) -> Any:
11        """Deserialize a model from bytes."""
12        ...
13
14    @abstractmethod
15    def info(self, model, /) -> dict | None:
16        """Extract structural information as JSON."""
17        ...
18
19    @abstractmethod

```

```

20     def version(self) -> str:
21         """Return the installed framework version."""
22         ...

```

A factory class (`ModelSerializer`) selects the appropriate serializer by inspecting the model object’s module hierarchy. A model whose class resides under `sklearn.*` dispatches to the scikit-learn serializer, one under `torch.*` to the PyTorch serializer, and so on, so a live model object never needs its framework declared. The one case that does is a directory path, which carries no class to inspect: passing one to `dump()` without a `lib` argument raises `ValueError`.

14.1 Scikit-learn

Operation	Implementation
<code>dump()</code>	Serializes the model with <code>pickle.dumps()</code> and encodes the result as Base64. Pickle is the standard serialization format for scikit-learn (Pedregosa et al., 2011) and supports the full range of estimator types, including pipelines and custom transformers.
<code>load()</code>	Decodes the Base64 string and deserializes with <code>pickle.loads()</code> .
<code>info()</code>	Calls <code>model.get_params()</code> and recursively sanitizes the result into a JSON-compatible dictionary (non-finite floats become <code>null</code>). Nested estimators (e.g. the base estimator in a <code>BaggingClassifier</code>) are serialized as their class names with parameters.

14.2 TensorFlow and Keras

Operation	Implementation
<code>dump()</code>	Saves the model to a temporary file in the native <code>.keras</code> format using the standalone Keras 3 package’s <code>keras.models.save_model()</code> , reads the file into memory, and encodes as Base64. The <code>.keras</code> format (Keras Team, 2023) captures the model architecture, weights, optimizer state, and compilation configuration in a single file.
<code>load()</code>	Writes the Base64-decoded bytes to a temporary file and loads with <code>keras.models.load_model(path)</code> .
<code>info()</code>	Calls <code>model.to_json()</code> and parses the JSON string into a dictionary. The result is a complete description of the model’s layer structure, including layer types, output shapes, activation functions, and connections.

14.3 PyTorch

Operation	Implementation
<code>dump()</code>	Exports the model to an <code>ExportedProgram</code> (PyTorch Team, 2024) via <code>torch.export.export(obj, example_inputs, dynamic_shapes)</code> , then serializes the exported graph to a byte buffer with <code>torch.export.save()</code> . The result is Base64-encoded. The export API captures the model’s computation graph with explicit dynamic-shape annotations, producing a portable, optimizable artifact that is independent of the Python class definition. If <code>example_inputs</code> is not provided, the serializer infers a suitable shape by inspecting the first <code>torch.nn.Linear</code> layer in the module.
<code>load()</code>	Decodes the Base64 string and loads the exported program with <code>torch.export.load(BytesIO(data))</code> . The loaded module can run inference without the original class definition or training code.
<code>info()</code>	Extracts the module list (as string representations), named parameters (as string representations, which include shape and dtype), and the full state dictionary with every tensor converted to a nested list. The result is self-contained, at the cost of size for large models, and requires no forward pass execution.

The `torch.export` API replaces the earlier TorchScript approach and offers several advantages for deployment: the exported graph preserves dynamic shapes (e.g. variable batch sizes) through explicit `Dim` annotations rather than trace-time heuristics, it supports the full Python operator set without the restrictions imposed by TorchScript’s subset compiler, and it integrates with PyTorch’s ahead-of-time compilation pipeline for further optimization at load time. The dynamic shapes are declared via a dictionary mapping input names to dimension constraints:

```

1 from flama.serialize import dump
2
3 dump(
4     model,
5     path="classifier.flm",
6     example_inputs=(torch.randn(2, 768),),
7     dynamic_shapes={"x": {0: torch.export.Dim("batch",
8                                     min=1)}}),
9 )

```

14.4 HuggingFace Transformers

The Transformers serializer follows a fundamentally different strategy from the other three: rather than a flat binary blob, it packages a whole directory, matching protocol version 2’s `bundle` model kind (Section 13).

Operation	Implementation
<code>dump()</code>	Accepts either a directory of pretrained model files or a live <code>transformers.Pipeline</code> (whose weights are first written out with <code>save_pretrained()</code> to a temporary directory). Either way, the directory is packed into an uncompressed tar archive in memory using the framework’s own Rust-accelerated tar routine, since the on-disk files (safetensors shards, tokenizer, configuration) are already in efficient formats and gain little from re-compression.
<code>load()</code>	Takes the path to the bundle already extracted to disk (raw bytes are rejected; Transformers reads a snapshot directory, not a byte stream) and calls <code>transformers.pipeline(task=task, model=str(path), **kwargs)</code> , returning a ready-to-use <code>transformers.Pipeline</code> .
<code>info()</code>	Reads back <code>pipeline.model.config.to_dict()</code> (the model architecture and hyperparameters), <code>pipeline.task</code> , and <code>pipeline.model.name_or_path</code> .

Capability detection is correspondingly more involved than for the other three frameworks, which always report an empty capability set. The Transformers serializer inspects the bundle on disk: `vision_config/audio_config` blocks in `config.json` (corroborated against the actual tensor names in the safetensors header, since a checkpoint can ship a multimodal config without the corresponding tower weights) drive image and audio support, and tool and reasoning support are detected by rendering the tokenizer’s chat template against sentinel inputs and checking whether the rendered output reflects them.

15 Model wrappers

Between the serialized model on disk and the HTTP endpoint that serves predictions, there are two cooperating layers: a **model wrapper**, engine-agnostic and shared by every framework, and a **backend**, one concrete class per framework, that the wrapper delegates the actual inference call to. The split exists because the two concerns change for different reasons: the wrapper owns lazy deserialization, metadata access, and the DI/HTTP-facing surface, none of which differ across frameworks, while

inference semantics do differ across frameworks (a scikit-learn estimator expects a NumPy array; a PyTorch module expects a `torch.Tensor`; a Transformers pipeline tokenizes internally) and are confined to the backend.

The wrapper base class defines lazy access to the model’s metadata, bundled artifacts, and backend, plus `inspect()`:

```

1 class BaseModel(Generic[B]):
2     def __init__(self, backend: B | None = None,
3                 meta: Metadata | None = None,
4                 artifacts: Artifacts | None = None, *,
5                 name: str | None = None,
6                 path: Path | None = None,
7                 autoload: bool = False):
8         ...
9
10    def inspect(self) -> dict:
11        return {"meta": self.meta.to_dict(),
12                "manifest": list(self.manifest)}
```

`manifest` is the list of bundled artifact *names*, read cheaply from the FLM header; the artifacts themselves are only extracted to disk once the model is actually loaded. `MLModel(BaseModel[MLBackend])` is the single concrete wrapper for every predictive framework; it adds `predict(x)` and `stream(x)`, both implemented once and delegated straight to `self.backend.predict(x)`. There is no per-framework subclass of the wrapper: the framework-specific code lives entirely in the backend that `MLBackend.from_model_artifact()` selects at load time, based on `Metadata.framework.lib`.

15.1 Framework-specific backends

Each framework provides a concrete `MLBackend` subclass. Its `predict(x)` method translates the framework-neutral input (a list of lists, received as JSON from the client) into the framework’s native input format, runs inference, and converts the output back to a JSON-serializable list:

Framework	<code>predict(x)</code> implementation
scikit-learn	<code>self.model.predict(x).tolist()</code>
TensorFlow	<code>self.model.predict(np.array(x)).tolist()</code>
PyTorch	<code>self.model(torch.Tensor(x)).tolist()</code>
Transformers	<code>self.model(x)</code> , delegating entirely to a <code>transformers.Pipeline</code> ’s own <code>tokenize/infer/decode</code> cycle.

In every case, `self.model` is the object the framework’s serializer produced at load time (an estimator, a Keras model, an exported graph module, or a `transformers.Pipeline`), and each backend raises `FrameworkNotInstalled` up front if the underlying library is not importable, rather than failing with an opaque `ImportError` mid-request.

16 Three levels of model integration

FLAMA provides three progressively more automated levels for integrating a machine-learning model into a web application. Each level builds on the one below it, and the developer chooses the level that matches the degree of customization required. At one extreme, the developer writes the endpoint

code and merely uses the framework for model loading and lifecycle management. At the other extreme, the developer provides only a model file path and a name, and the framework generates the complete endpoint infrastructure automatically.

16.1 Level 1: Model components

At the lowest level, a model is loaded as a dependency-injection component. The `ModelComponentBuilder` factory reads the metadata header from an FLM file to determine the artifact family, then produces a `ModelComponent` that makes the model available for injection into any handler:

```
1 from flama.models.components import ModelComponentBuilder
2
3 # Build a lazy component from a .flm file
4 component = ModelComponentBuilder.build("classifier.flm")
```

The builder performs the following steps:

1. Read the FLM file's metadata header (a cheap operation that does not deserialize the model body).
2. Determine the artifact family ("ml" or "llm") from the metadata.
3. For ML artifacts, create a fresh, per-instance subclass of `MLModel`; for LLM artifacts, of `LLMModel`. The subclass carries no extra behaviour of its own: its only purpose is to give this particular registered model its own type.
4. Wrap it in a `ModelComponent` subclass whose `resolve()` method returns that instance, with its return type annotation set to the freshly created subclass. Because the injector keys components by the return annotation of `resolve()`, this is what lets two different registered models, each with its own dynamically created type, resolve to two different handler parameters without colliding.

The actual model body is not deserialized at build time. Heavy loading (backend initialization, weight deserialization) is deferred to `component.startup()`, so the server port binds before model loading begins.

A handler can then receive the model via dependency injection. Since the component's type is generated at build time rather than imported, the handler is typed against `component.get_model_type()` and the component is registered explicitly, both on the application and on its startup event:

```
1 import typing as t
2 from flama import Flama, schemas
3
4 component = ModelComponentBuilder.build("classifier.flm")
5 Model = component.get_model_type()
6
7 app = Flama(events={"startup": [component.startup]})
8 app.add_component(component)
9
10 @app.post("/predict")
11 async def predict(
12     model: Model,
```



```

13     data: t.Annotated[schemas.Schema, schemas.SchemaMetadata(
14         PredictInput)],
15 ) -> t.Annotated[schemas.Schema, schemas.SchemaMetadata(PredictOutput)
16     ]:
17     result = model.predict(data["input"])
18     return {"output": result}

```

This level is appropriate when the developer needs full control over the endpoint's URL, methods, request/response schemas, and error handling, but still wants the model loading and lifecycle managed by the framework.

16.2 Level 2: The `add_model()` API

The `ModelsModule` provides a higher-level API that creates a complete model resource from a single method call:

```

1 app = Flama()
2
3 app.models.add_model(
4     path="/classifier",
5     model="classifier.flm",
6     name="classifier",
7 )

```

This call generates two endpoints:

Method	Path	Behaviour
GET	/classifier/	Returns the model's metadata (framework, version, class, hyperparameters, metrics, artifact list).
POST	/classifier/predict/	Accepts <code>PredictInput</code> (<code>{"input": [...]}</code>), runs inference, and returns <code>PredictOutput</code> (<code>{"output": [...]}</code>).

16.3 Level 3: Model resources

The most declarative level uses a class-based resource with a metaclass that generates all the infrastructure automatically:

```

1 from flama.models import MLResource
2
3 class SentimentClassifier(MLResource):
4     name = "sentiment"
5     verbose_name = "Sentiment Classifier"
6     model_path = "models/sentiment.flm"

```

Like a `CRUDResource`, the class declaration alone does not attach it to the application; it still has to be registered through the models module:

```

1 app.models.add_model_resource(path="/sentiment", resource=
    SentimentClassifier)

```

The `MLResourceType` metaclass performs the following operations when the class is defined:

1. Loads the model component from `model_path` using the same builder as Level 1.
2. Stores the component, the model wrapper, and the model type in the class's internal namespace.
3. Calls three mixins (`InspectMixin`, `PredictMixin`, `StreamMixin`) to generate endpoint methods. Each mixin's `_add_*`() method creates a handler function on the class with the correct signature and type annotations.

This level is appropriate for applications that deploy multiple models and want each model to be a self-contained, reusable unit with its own configuration.

17 Large language model serving

Everything so far has concerned *predictive* models: the binary format, the serializers, and the three levels of integration that put a trained classifier behind a REST endpoint. The *generative* case reuses that architecture but asks more of it, since a language model has to stream, has to speak several wire protocols at once, and has to run on whatever accelerator the host happens to have.

Those are the three requirements that separate it from the predictive case:

1. **Streaming output.** Language model inference produces tokens incrementally. Users expect to see output appear as it is generated, not after the entire sequence has been computed. The serving layer must therefore deliver partial results via Server-Sent Events or Newline-Delimited JSON as the model generates.
2. **Multi-protocol compatibility.** The generative AI ecosystem has converged on several incompatible wire protocols (OpenAI, Anthropic, Ollama). Practitioners integrate with these protocols using existing client SDKs and tooling. The serving layer must expose the same physical model through all major protocols without duplicating the inference pipeline.
3. **Hardware-aware backends.** LLM inference is compute-bound and benefits from hardware-specific optimization (PagedAttention on CUDA, Metal acceleration on Apple Silicon). The serving layer must select the highest-performance backend available at runtime without requiring the user to write backend-specific code.

17.1 Architecture overview

The LLM serving subsystem is organized as a five-layer pipeline:

Wire dialect parses incoming request JSON into canonical transport objects (messages, tools) and renders outgoing events into protocol-specific streaming frames or buffered envelopes. Four dialects are provided: OpenAI, Anthropic, Ollama, and Native.

Transport defines the canonical request and response model, independent of any wire protocol. A request is a sequence of typed messages with multimodal content parts; a response is a stream of typed events (start, text, tool call, trace, stop).

Engine bridges the transport layer and the backend. It converts canonical messages into tokenized `EngineInput` (token IDs plus decoded media), invokes the backend's generation routine, and produces a stream of `EngineDelta` objects carrying incremental text, token counts, and finish reasons.

Codec/Decoder transforms the raw **EngineDelta** stream into canonical output events. The **LLMCodec** implements a finite-state machine that recognizes reasoning channels (think tags, channel markers) and tool-call bodies in the generated text, emitting structured **TextEvent**, **ToolEvent**, and **TraceEvent** objects.

Backend is the hardware-bound inference engine. Two backends are supported: **vLLM** (Linux/CUDA) and **MLX** (Apple Silicon). The backend is selected at model load time by probing which library is importable.

A typical end-to-end inference flow proceeds as follows:

1. An HTTP request arrives at a dialect-specific endpoint (e.g. `/v1/chat/completions`).
2. The dialect's *parser* converts the wire-format JSON into canonical **Message** and **Tool** objects.
3. A **Shape** (raw, chat, or conversation) renders the messages into **EngineInput** by applying the backend's chat template and tokenizer.
4. The backend generates tokens, yielding **EngineDelta** objects incrementally.
5. The **LLMCodec** decodes deltas into canonical **Event** objects.
6. An **EventBuffer** feeds events into the dialect's *renderer*, which produces SSE or NDJSON frames.
7. For buffered (non-streaming) requests, the dialect's *assembler* coalesces all events into a single response envelope.

17.2 Backends

A backend is responsible for loading a model into memory, applying a chat template to format messages, tokenizing input, and generating tokens. The backend abstraction is minimal: any object that exposes `encode()`, `chat_template()`, `prepare_input()`, and `generate()` can serve as a backend. Two implementations are provided:

vLLM ([Kwon et al., 2023](#)) is a high-throughput inference engine for Linux systems with NVIDIA GPUs. It implements PagedAttention for efficient KV-cache management, continuous batching for maximizing GPU utilization, and tensor parallelism for multi-GPU deployments. **vLLM** is the preferred backend for production deployments where throughput and latency are critical.

MLX ([Hannun et al., 2023](#)) is Apple's framework for machine learning on Apple Silicon. It provides Metal-accelerated tensor operations with a NumPy-compatible API and unified memory that eliminates data transfers between CPU and GPU. The **MLX** backend uses `mlx-lm` for text generation and `mlx-vlm` for vision-language models. It is the preferred backend on macOS systems.

Both backends share the **TransformerLLMBackend** abstract base class, which delegates chat-template rendering to the Hugging Face `tokenizer` and `AutoProcessor` stack. This design means that any model checkpoint compatible with the Hugging Face model format can be loaded by either backend; the runtime engine is what differs, not the model loading or template application logic.

Backend selection is automatic: at model load time, the framework probes which libraries are importable and selects the first available option. On a Linux system with CUDA, **vLLM** is used; on macOS with Apple Silicon, **MLX** is used. If neither is available, the framework raises a clear

error indicating that one of the two runtimes must be installed. The selection is transparent to the application: the same `add_model()` call works identically regardless of the underlying backend.

17.3 Transport layer

The transport layer defines the canonical data model that sits between the wire protocols and the engine. It comprises three sub-layers: input shapes, messages, and output events.

17.3.1 Input shapes

A **Shape** determines how a client's messages are formatted before they reach the tokenizer. Three shapes are supported:

Raw sends the prompt text directly to the tokenizer without any formatting. This is suitable for completions-style inference where the client provides the exact token sequence.

Chat wraps the input into a two-message sequence (optional system message plus user message) and applies the model's chat template. This is the default shape for single-turn interactions.

Conversation passes the full message history (system, user, assistant, tool results) through the chat template. This enables multi-turn dialogues where the model has access to the entire conversation context.

17.3.2 Messages and content parts

The **Message** hierarchy represents a single conversational turn. Each message has a role (**system**, **user**, **assistant**, or **tool**) and a sequence of typed content parts:

- **Text**: plain text content.
- **Image**: a base64-encoded image or a URL, with MIME type and optional detail level.
- **Audio**: a base64-encoded audio segment with sample rate and format metadata.
- **Tool call**: a function invocation request from the assistant, with a call ID, function name, and JSON arguments.
- **Tool result**: the return value of a tool call, associated with the original call ID.

This multimodal message representation is the canonical form into which every wire protocol's request is parsed. The dialect parsers handle the translation from protocol-specific formats (OpenAI's content-part arrays, Anthropic's top-level system field, Ollama's sibling **images** array) into this uniform representation.

17.3.3 Output events

The response from an LLM is represented as a stream of typed events:

StartEvent signals the beginning of a generation. Carries the model name and the generation configuration.

TextEvent delivers a fragment of generated text. In streaming mode, one **TextEvent** is emitted per decoded token or token group.

ToolEvent delivers a complete tool-call request extracted from the generated text. Carries the function name, JSON arguments, and an optional call ID. Tool calls are emitted atomically (the full call body is accumulated before emission) to ensure well-formed JSON.

TraceEvent delivers content from a secondary output channel (e.g. reasoning traces, thinking tokens, analysis steps). Trace events are emitted when the codec detects channel markers in the generated text.

StopEvent signals the end of generation. Carries the stop reason (end-of-sequence, length limit, tool call, or content filter), token usage statistics, and optional metadata.

17.4 Codec and decoder pipeline

The **LLMCodec** is the bridge between the backend’s raw token stream and the canonical event model. Its core responsibility is *structured output recognition*: detecting reasoning channels and tool calls in the generated text and emitting appropriate events rather than treating all output as undifferentiated text.

17.4.1 Decoder detection

Different model families use different conventions for structured output. Some models wrap reasoning in `<think>...</think>` tags; others use `<|begin_of_thought|>` markers; still others emit tool calls as JSON objects preceded by special tokens. The **Decoder** is responsible for detecting which conventions a particular model uses.

Detection proceeds in three stages:

1. **Pinned configuration**: if the user explicitly specifies a channel scanner, tool scanner, or tool parser, those are used directly without auto-detection.
2. **Chat template analysis**: the decoder examines the model’s chat template for known marker patterns and selects scanners accordingly.
3. **Preflight probing**: if template analysis is inconclusive, a short preflight generation is performed and the output is analysed for structural patterns.

The decoder comprises three pluggable components:

Channel scanner recognizes the start and end of secondary output channels. Built-in scanners support `think` tags, generic `channel` markers, and passthrough (no channels).

Tool scanner recognizes the start of a tool-call body in the generated text. It detects special tokens (e.g. `<tool_call>`) or JSON-object openings that signal a function invocation.

Tool parser extracts the function name and arguments from the scanned tool-call body. Built-in parsers support JSON objects, JSON arrays, named JSON sequences, and Python-style call notation.

17.4.2 Finite-state decoding

The codec maintains a three-state finite-state machine that processes each **EngineDelta**:

- **Outside:** the default state. Text is emitted as `TextEvent`. If the channel scanner detects a channel opening, the state transitions to *channel*. If the tool scanner detects a tool-call start, the state transitions to *tool*.
- **Channel:** text is accumulated as trace content and emitted as `TraceEvent`. When the channel scanner detects the closing marker, the state returns to *outside*.
- **Tool:** text is accumulated until the tool body is complete (balanced braces or explicit end marker). The tool parser then extracts the function name and arguments, and a `ToolEvent` is emitted atomically. The state returns to *outside*.

This architecture ensures that the downstream dialect renderers receive clean, typed events regardless of how the model formats its output. A model that emits tool calls as raw JSON in a `<tool_call>` block and a model that uses Python-style `function(args)` notation both produce identical `ToolEvent` objects.

17.5 Wire dialects

A **Dialect** is the complete adapter between a wire protocol and the canonical transport model. It comprises three components:

Parser converts wire-format request JSON into canonical `Message` and `Tool` objects. Each protocol has different conventions for representing messages (OpenAI uses content-part arrays; Anthropic uses a top-level system field; Ollama uses sibling image arrays), and the parser normalizes these into the common representation.

Renderer converts canonical output events into streaming wire frames (SSE or NDJSON). The renderer is invoked incrementally as events arrive and produces one or more protocol-specific frames per event.

Assembler converts a complete sequence of canonical events into a buffered response envelope. The assembler is used for non-streaming requests where the client expects a single JSON response rather than a stream.

Four dialects are provided:

17.5.1 OpenAI dialect

The OpenAI dialect implements the Chat Completions, Completions, Responses, and Models APIs. It is mounted at `/v1/chat/completions`, `/v1/completions`, `/v1/responses`, and `/v1/models`. Streaming responses use SSE with `chat.completion.chunk` or `text_completion` events. Buffered responses return `chat.completion` or `response` envelopes.

This dialect enables any client library that targets the OpenAI API (the OpenAI Python SDK, LangChain, LlamaIndex, or any HTTP client) to interact with a FLAMA-served model without modification.

17.5.2 Anthropic dialect

The Anthropic dialect implements the Messages and Models APIs, mounted at `/v1/messages` and `/v1/models`. Streaming uses SSE with Anthropic-specific event types (`message_start`, `content_`

`block_start`, `content_block_delta`, `message_delta`, `message_stop`). The parser handles Anthropic's conventions: top-level `system` field, `thinking` content blocks for reasoning traces, and `tool_use` blocks with explicit IDs.

17.5.3 Ollama dialect

The Ollama dialect implements the Chat, Generate, Tags, Show, and Version APIs, mounted at `/api/chat`, `/api/generate`, `/api/tags`, `/api/show`, and `/api/version`. Unlike the other dialects, Ollama uses Newline-Delimited JSON (NDJSON) for streaming rather than SSE. The parser normalizes Ollama's sibling `images` arrays into structured multimodal content parts.

17.5.4 Native dialect

The native dialect is FLAMA's own protocol, designed for maximum fidelity to the internal event model. Unlike the other three, it takes no URL prefix of its own, so its routes sit directly under the model's mount point: `/` configures the model's generation parameters, `/query/` returns a buffered response, `/stream/` opens a live stream, `/stream/{stream_id}/` replays a past one, and `/chat/` serves the chatbot interface.

The streaming format uses SSE with sequential event IDs that enable client reconnection via the `Last-Event-ID` header. The buffered `/query/` response is not assembled by the dialect the way the other three assemble theirs: the native `Assembler` raises `NotImplementedError`, because the wire format has no buffered envelope of its own. Instead the handler coalesces the event stream inline and returns the canonical event model directly, as an envelope carrying `id`, `created`, `stop_reason`, token counts, and the channel-tagged `blocks` (one per contiguous run on the same channel).

The native dialect also mounts a built-in chatbot web interface at `/chat/` (Section 17.8).

17.6 Serving configuration

An LLM model is added to an application with the same `add_model()` API used for predictive models, with the addition of a `serving` parameter that declares which dialects to expose:

```
1 from flama import Flama
2
3 app = Flama()
4
5 app.models.add_model(
6     path="/llm/",
7     model="assistant.flm",
8     name="assistant",
9     serving=("native", "openai", "anthropic", "ollama"),
10    params={"temperature": 0.7, "max_tokens": 512},
11 )
```

This single call generates the complete set of endpoints for all four dialects, mounted under the specified path prefix. The resulting URL structure is:

Dialect	Endpoint	Format
Native	/llm/query/	JSON (buffered blocks)
Native	/llm/stream/	SSE (canonical events)
Native	/llm/chat/	SSE + chatbot UI
OpenAI	/llm/openai/v1/chat/completions	SSE / JSON
Anthropic	/llm/anthropic/v1/messages	SSE / JSON
Ollama	/llm/ollama/api/chat	NDJSON / JSON

Generation parameters (`temperature`, `top_p`, `max_tokens`, etc.) can be set at the serving level as defaults and overridden per-request by the client through protocol-specific fields.

17.7 Stream persistence and replay

The native dialect supports stream persistence, enabling clients to replay past generations and resume interrupted streams. The `StreamsBackend` abstraction provides durable event storage, keyed throughout by a (`model`, `stream id`) pair: `append()` adds one event to a stream, `read()` returns a half-open range of it, `pop()` reads a range and drops it, `discard()` removes a stream outright, and `length()` reports how many events each stream currently holds. Backends are opened and closed with the application through `aopen()/aclose()`.

Two implementations ship with the framework. The default `FileStreamsBackend` writes each stream to `<root>/<model>/<stream_id>.jsonl`, one self-contained JSON event per line, and deliberately leaves those logs in place after shutdown so they remain available for inspection; `InMemoryStreamsBackend` keeps the same interface without touching disk. Range reads are what make replay cheap: a client reconnecting with a `Last-Event-ID` header is served the stored events from that sequence number onward, followed by live generation for whatever has not been produced yet. A `CleanupTask` bounds the cost of retention, evicting streams by age and by aggregate disk usage.

17.8 Chatbot template

The native dialect includes a built-in chatbot web interface that provides a complete conversational UI without any frontend code. The interface is a self-contained HTML page rendered from a compiled template and served at the `/chat/` endpoint.

The chatbot UI provides:

- Real-time token streaming over SSE with automatic reconnection.
- Markdown rendering for structured responses.
- $\text{\LaTeX}$ mathematics rendering via KaTeX for models that produce mathematical content.
- Mermaid diagram rendering for models that produce structured diagrams.
- Syntax highlighting for fenced code blocks.
- Conversation history management with multi-turn context.

The page is served by the native dialect's `/chat/` handler, which renders the packaged `chatbot/chat.html` template with the model's stream URL as its only context variable. The template ships with the framework rather than being a configuration point: an application that wants a different interface replaces the route with its own handler and consumes the same `/stream/` endpoint, which is the public contract the bundled page itself is written against.

18 The Model Context Protocol

The Model Context Protocol (MCP) ([Anthropic, 2026](#)) is an open standard for connecting AI models to external tools, data sources, and programmatic capabilities. As AI applications increasingly require interaction with external systems (databases, APIs, file systems, code execution environments), a standardized protocol for discovering and invoking these capabilities becomes necessary. FLAMA includes a first-class MCP module that allows any FLAMA application to act as an MCP server, exposing tools, resources, and prompts over a JSON-RPC 2.0 transport.

18.1 MCP server

An MCP server in FLAMA is a registry of three kinds of capabilities:

Tools are callable functions that perform actions: querying a database, calling an external API, running a computation, or controlling a device. Each tool has a name, a description, and an input schema generated from the handler's function signature.

Resources are read-only data sources identified by URIs. A resource handler returns the content of the resource (text, JSON, or binary) when requested by the client. Resources are used to expose static or dynamic data to the model.

Prompts are reusable prompt templates with parameters. The MCP client can list available prompts, retrieve a specific prompt with arguments filled in, and use the result as input to the language model.

```
1 from flama import Flama
2 from flama.mcp import MCPServer
3
4 mcp = MCPServer(
5     name="data-tools",
6     version="1.0.0",
7     instructions="Tools for querying the data warehouse.",
8 )
9
10 @mcp.tool(description="Run a SQL query")
11 def query_db(sql: str, limit: int = 100) -> str:
12     result = execute_query(sql, limit=limit)
13     return format_as_table(result)
14
15 @mcp.resource(uri="schema://tables",
16               description="List all database tables")
17 def list_tables() -> str:
18     return "\n".join(get_table_names())
19
20 @mcp.prompt(description="Generate a data analysis prompt")
21 def analyze(table: str, question: str) -> list:
22     return [
23         {"role": "system",
24          "content": "You are a data analyst."},
25         {"role": "user",
26          "content": f"Table: {table}\n{question}"},
27     ]
28
```

```

29 app = Flama()
30 app.mcp.add_server("/mcp", name="data-tools",
31                     server=mcp)

```

18.2 Stateless JSON-RPC transport

The MCP module communicates over HTTP using the JSON-RPC 2.0 protocol ([JSON-RPC Working Group, 2013](#)). Each MCP server is mounted as a POST-only HTTP endpoint (hidden from the OpenAPI schema) that accepts JSON-RPC requests and returns JSON-RPC responses.

FLAMA implements the current stateless revision of the MCP specification, in which every request is self-contained: there is no `initialize/initialized` session handshake. Instead, the client’s identity, capabilities, and requested protocol version travel in the `_meta` field of each JSON-RPC request. The server validates the protocol version, extracts client capabilities, and responds accordingly.

Three routing headers accompany every request to enable efficient HTTP-level dispatch:

- **Mcp-Method**: the JSON-RPC method name (e.g. `tools/call`).
- **Mcp-Name**: the target name (tool name, resource URI, or prompt name), extracted from the body’s `params` field.
- **MCP-Protocol-Version**: the protocol revision requested by the client.

The server validates that these headers are consistent with the request body before dispatching the call. Every response carries the **MCP-Protocol-Version** header to confirm the negotiated version.

The endpoint handles the following methods:

JSON-RPC method	Behaviour
<code>server/discover</code>	Returns the server’s capabilities (tools, resources, prompts), supported extensions, server info, and optional instructions. Replaces the earlier session-based <code>initialize</code> handshake.
<code>ping</code>	Returns an empty object. Used for health checks.
<code>tools/list</code>	Returns the list of registered tools with names, descriptions, input schemas (JSON Schema 2020-12), and output schemas.
<code>tools/call</code>	Invokes a named tool with arguments and returns the result as a content block. Supports both immediate and task-based (asynchronous) execution.
<code>resources/list</code>	Returns registered resources with URIs and MIME types.
<code>resources/read</code>	Returns the content of a resource identified by URI.
<code>prompts/list</code>	Returns registered prompts with names and descriptions.
<code>prompts/get</code>	Returns a prompt’s messages with arguments filled in.
<code>tasks/get</code>	Retrieves the current status of an asynchronous task.
<code>tasks/cancel</code>	Cancels a running task.

Every JSON-RPC response includes a `_meta.dev.flama` field with FLAMA’s branding information, allowing MCP clients to identify the server framework.

18.2.1 Tool schemas

Tool input and output schemas are emitted as self-contained JSON Schema 2020-12 documents with local `$defs` references. Input schemas are generated automatically from the tool handler’s function signature (parameter names, types, and defaults). Output schemas are generated from the return type annotation when present; tools without a return annotation advertise no output schema.

```

1 @mcp.tool(description="Add two numbers")
2 def add(a: float, b: float) -> float:
3     return a + b

```

This tool produces an input schema with two required `number` properties and an output schema declaring a `number` result, both as standalone JSON Schema documents that clients can validate against independently.

18.2.2 Trace context

The MCP module extracts W3C Trace Context ([W3C, 2021](#)) fields from the request's `_meta: traceparent`, `tracestate`, and `baggage`. These are exposed as an injectable `TraceContext` component that tool handlers can use to propagate distributed traces to downstream services.

18.3 Tasks extension

Some tool invocations are inherently long-running: training a model, executing a complex query, or waiting for an external approval. The MCP Tasks extension allows such tools to return immediately with a task identifier that the client can poll for completion.

A tool opts into task-based execution by declaring `task=True` at registration time:

```

1 @mcp.tool(description="Train a model", task=True)
2 async def train(dataset: str, epochs: int = 10) -> str:
3     result = await run_training(dataset, epochs)
4     return f"Model trained: accuracy {result.accuracy}"

```

When a client whose capabilities include the Tasks extension invokes this tool, the server returns a task object immediately with status `running`. The framework executes the handler in the background and updates the task's status (`running` → `completed` or `failed`) as the handler progresses. The client retrieves the result via `tasks/get`.

The task store persists task state across requests, and clients can cancel running tasks via `tasks/cancel`.

18.4 Elicitation extension

Elicitation enables a tool handler to request additional input from the user mid-execution. This is useful for tools that need confirmation, disambiguation, or iterative refinement before completing their work.

A tool signals that it needs more input by returning an `Elicit` object:

```

1 from flama.mcp import Elicit, Elicitation
2
3 @mcp.tool(description="Delete records")
4 def delete_records(table: str, where: str,
5                   elicitation: Elicitation) -> str:
6     if "confirm" not in elicitation:
7         return Elicit.require(
8             message=f"Delete from {table} where {where}?",
9             schema={"confirm": {"type": "boolean"}},

```

```

10         )
11
12         if elicitation.get("confirm"):
13             count = execute_delete(table, where)
14             return f"Deleted {count} records."
15         return "Cancelled."

```

When the handler returns an `Elicit`, the server responds with `resultType: inputRequired` and a `requestState` token that encodes the continuation context. The client presents the elicitation prompt to the user, collects the response, and sends a follow-up `tools/call` with the user's answers and the `requestState` token. The framework deserializes the state, injects the collected answers as the `Elicitation` parameter, and re-executes the handler.

This mechanism is fully stateless: the continuation state is serialized into the response and returned by the client on the next request, with no server-side session required.

18.5 Application templates extension

The MCP Apps extension allows servers to expose prefetchable UI templates that clients can render inline. An application template is a named HTML or Markdown document that the client can retrieve and display as part of a tool's result.

```

1  @mcp.app_template(
2      "ui://chart",
3      name="chart",
4      description="Render a data chart",
5      mime_type="text/html",
6  )
7  def chart_template(data: str) -> str:
8      return render_chart_html(data)

```

Templates are discovered through `resources/templates/list` and read through `resources/read`. The server advertises the Apps extension in its capabilities when at least one template is registered.

18.6 Multiple servers

A single FLAMA application can host multiple MCP servers, each mounted at a different path and exposing a different set of capabilities:

```

1  app.mcp.add_server("/mcp/data", name="data-tools",
2                      server=data_mcp)
3  app.mcp.add_server("/mcp/admin", name="admin-tools",
4                      server=admin_mcp)

```

Tools, resources, and prompts can also be registered directly on the application using the `MCPModule`'s decorator API, targeting a specific server by name:

```

1  @app.mcp.tool(description="Health check",
2                mcp="admin-tools")
3  def health() -> str:
4      return "OK"

```

If only one MCP server is registered, the `mcp` parameter can be omitted and the tool is added to the sole server automatically.

Part IV

Operations and Tooling

19 Command-line interface

Six commands, built on Click ([Ronacher, 2014](#)), cover the operational lifecycle: running an application from an import path, serving a model file with no application code at all, driving multi-model deployments from a configuration file, fetching and packaging models from a remote repository, working with a model offline from the terminal, and migrating a codebase across a major version. Two audiences use them, and the second matters more for the design than the first. Developers reach for the CLI interactively, but deployment scripts, CI pipelines, and container orchestrators reach for it unattended, which is why every option can also be set from a `FLAMA_*` environment variable. Each command is set out below with its options and a worked example.

19.1 Overview

The CLI is invoked through the `flama` entry point and organizes its functionality into six top-level commands:

```
$ flama --help
Usage: flama [OPTIONS] COMMAND [ARGS]...

    Fire up your models with Flama

Options:
  --version  Check the version of your locally installed Flama
  --help     Get help about how to use Flama CLI

Commands:
  get      Download and package a model as .flm.
  model    Interact with an ML model without server.
  run      Run a Flama Application based on a route.
  serve    Serve an ML model file within a Flama Application.
  start    Start a Flama Application based on a config file.
  upgrade  Upgrade a Flama codebase to a newer major version.
```

All options support environment variable binding through the `FLAMA_*` prefix (e.g. `FLAMA_APP=mymodule:app flama run`). This makes the CLI fully compatible with containerized deployments where configuration is injected through environment variables rather than command-line arguments.

19.2 flama run

The `run` command starts a Uvicorn ([Christie, 2017](#)) server hosting a user-defined FLAMA application identified by its Python import path:

```
$ flama run myapp:app --server-host 0.0.0.0 --server-port 8000
```

The argument `myapp:app` follows the standard ASGI convention: the module path on the left of the colon and the application variable name on the right. The command accepts the full set of Uvicorn server options, prefixed with `--server-`, including:

Option	Default	Type	Purpose
<code>--server-host</code>	127.0.0.1	str	Bind address
<code>--server-port</code>	8000	int	Bind port
<code>--server-reload</code>	off	flag	Auto-reload on code changes
<code>--server-workers</code>	1	int	Number of worker processes
<code>--server-loop</code>	auto	choice	Event loop implementation
<code>--server-http</code>	auto	choice	HTTP protocol implementation
<code>--server-ws</code>	auto	choice	WebSocket implementation
<code>--server-log-level</code>	info	choice	Logging verbosity
<code>--server-ssl-certfile</code>	—	path	SSL certificate file
<code>--server-ssl-keyfile</code>	—	path	SSL private key file

The `run` command is the appropriate choice when the developer has written a complete FLAMA application with custom routes, middleware, and configuration.

19.3 flama serve

The `serve` command is a one-liner for deploying one or more models behind a REST API with no application code. It accepts one or more `--model` specifications and generates a FLAMA application on the fly:

```
$ flama serve --model classifier.flm
```

This command:

1. Reads the model file and detects the artifact family from its metadata.
2. Generates a temporary Python module from a Jinja2 template (`app.py.j2`) that creates a FLAMA application and registers each model via `app.models.add_model()`.
3. Starts a Uvicorn server hosting the generated application.

The generated application exposes each model at its specified URL prefix, plus an OpenAPI schema at `/schema/` and a Swagger UI at `/docs/`.

Application-level options customize the generated application:

Option	Default	Env var	Purpose
<code>--app-title</code>	Flama	APP_TITLE	Application name
<code>--app-version</code>	0.1.0	APP_VERSION	Application version
<code>--app-description</code>	(default)	APP_DESCRIPTION	OpenAPI description
<code>--app-debug</code>	off	APP_DEBUG	Enable debug mode
<code>--app-schema</code>	<code>/schema/</code>	APP_SCHEMA	OpenAPI schema route
<code>--app-docs</code>	<code>/docs/</code>	APP_DOCS	Swagger UI route

The `--model` option accepts three forms. The simplest is a bare path; the full form uses comma-separated key=value pairs; and the file form loads a complete specification from JSON, YAML, or TOML:

```
# Bare path (defaults: url=/, name=model)
$ flama serve --model classifier.flm

# Full form with explicit keys
$ flama serve --model file=classifier.flm,url=/sentiment,name=v2

# Load spec from file
$ flama serve --model @spec.json
```

The model specification supports the following keys:

Key	Default	Applies to	Purpose
file	(required)	all	Path to the .flm file
url	/	all	URL prefix for the model's endpoints
name	model	all	Model name
serving	(auto)	LLM	Colon-separated dialect list (native, openai, anthropic, ollama)
params	(none)	LLM	Colon-separated key=value generation parameters
channel_scanner	(auto)	LLM	Override channel detection
tool_scanner	(auto)	LLM	Override tool-call detection
tool_parser	(auto)	LLM	Override tool-call parsing

A complete LLM deployment with multiple dialects:

```
$ flama serve \
  --model file=assistant.flm,url=/llm,name=assistant,\
  serving=native:openai:anthropic,\
  params=temperature=0.7:max_tokens=4096 \
  --server-host 0.0.0.0 \
  --server-port 8000
```

This produces a single application with the native chatbot UI at `/llm/chat/`, OpenAI-compatible endpoints at `/llm/openai/v1/chat/completions`, and Anthropic-compatible endpoints at `/llm/anthropic/v1/messages`. Multiple `--model` flags can be passed to serve several models in the same application.

19.4 flama start

The `start` command is designed for multi-model deployments managed through a JSON configuration file. It reads a configuration file (default: `flama.json`) that specifies the application metadata, the list of models to serve, and the server options:

```
{
  "app": {
    "title": "ML Platform",
    "version": "2.0.0",
    "description": "Production model serving",
    "debug": false,
    "schema": "/schema/",
    "docs": "/docs/"
  }
}
```



```

"models": [
  {
    "url": "/sentiment",
    "path": "models/sentiment.flm",
    "name": "sentiment"
  },
  {
    "url": "/toxicity",
    "path": "models/toxicity.flm",
    "name": "toxicity"
  },
  {
    "url": "/assistant",
    "path": "models/assistant.flm",
    "name": "assistant",
    "serving": ["native", "openai"]
  }
],
"server": {
  "host": "0.0.0.0",
  "port": 8000,
  "workers": 4,
  "log_level": "info"
}
}

```

The command accepts a `--create-config` option that generates a template configuration file:

```

# Generate a minimal config with host and port only
$ flama start --create-config simple

# Generate a full config with all server options
$ flama start --create-config full

```

This is particularly useful for teams that manage model deployments through infrastructure-as-code workflows: the configuration file can be version-controlled, reviewed, and deployed through CI/CD pipelines.

19.5 flama get

The `get` command downloads a model from a remote source and packages it into the `.flm` format, ready for serving with `flama serve` or offline interaction with `flama model`. This command bridges the gap between model repositories and the FLAMA deployment pipeline.

```

$ flama get --source huggingface --family llm \
  Qwen/Qwen3-8B

```

The command downloads all model files concurrently (with configurable parallelism), packages them into a `.flm` file using protocol version 2, and records the artifact family in the manifest. The `--family` flag is required and determines how the model is dispatched at load time:

Option	Default	Type	Purpose
-source	(required)	choice	Model source provider (currently <code>huggingface</code>)
-family	(required)	choice	Artifact family (ml or llm). Drives runtime dispatch
-o, -output	(auto)	path	Output .flm path (default: <model-name>.flm)
-max-concurrent	8	int	Maximum parallel file downloads

Downloads employ bounded exponential backoff with jitter for transient failures (HTTP 429, 5xx, network errors), and each file is streamed to a temporary path and atomically renamed on success to prevent partial downloads from corrupting the output.

```
# Download a traditional ML model
$ flama get --source huggingface --family ml \
  scikit-learn/Fish-Weight

# Download an LLM with custom output path
$ flama get --source huggingface --family llm \
  -o assistant.flm \
  google/gemma-4-E2B-it
```

19.6 flama model

The `model` command group provides offline interaction with serialized models, without starting a server. It accepts a model path as a positional argument and exposes three subcommands: `inspect` (metadata display), `run` (one-shot inference), and `stream` (streaming inference). Both `run` and `stream` work for ML and LLM models alike; for LLM models, additional options control the transport shape, system instructions, generation parameters, and output channels.

The group-level options `--channel-scanner`, `--tool-scanner`, and `--tool-parser` apply to LLM artifacts only and override the automatic decoder detection described in Section 17.4.

19.6.1 flama model inspect

The `inspect` subcommand displays the metadata stored in the FLM file:

```
$ flama model classifier.flm inspect --pretty
```

```
{
  "meta": {
    "id": "classifier-v2",
    "timestamp": "2025-01-15T10:30:00",
    "framework": {
      "lib": "sklearn",
      "version": "1.7.2"
    },
  },
  "model": {
    "obj": "RandomForestClassifier",
    "info": {
      "n_estimators": 100,
      "max_depth": 8,
      "criterion": "gini"
    },
  },
  "params": {
```

```

    "n_estimators": 100,
    "max_depth": 8
  },
  "metrics": {
    "accuracy": 0.947,
    "f1": 0.932
  }
},
"extra": {
  "dataset": "prod-2024-q3",
  "author": "team-ml"
}
},
"artifacts": {}
}

```

This subcommand is useful for verifying that a model was serialized correctly, checking its training metrics before deployment, and inventorying models in a model registry.

19.6.2 flama model run

The `run` subcommand performs one-shot inference without a server. For ML models, the input is a JSON array of feature vectors and the output is a JSON array of predictions. For LLM models, the input is a prompt and the output is the generated response:

```

# ML model: batch predictions from stdin
$ echo '[[5.1,3.5,1.4,0.2],[6.7,3.0,5.2,2.3]]' \
  | flama model classifier.flm run --pretty

# ML model: from file to file
$ flama model classifier.flm run \
  -i input.json -o output.json

# LLM model: single-turn generation
$ echo "Explain quantum entanglement briefly" \
  | flama model assistant.flm run --transport chat

# LLM model: with system instruction and params
$ echo "What is Python?" \
  | flama model assistant.flm run \
    --system "Be concise." \
    --param temperature=0.7

```

For LLM models, the `--transport` flag determines how the input is formatted before tokenization: `raw` sends the text directly, `chat` wraps it as a single user message with the model's chat template, and `conversation` expects a JSON array of messages. Generation parameters are passed as repeatable `--param key=value` flags.

The `--channel` flag includes secondary output channels (reasoning traces, analysis steps) in the output, which are otherwise suppressed by default. With one channel the output is plain text; with multiple channels each block is emitted as a JSON object with `channel` and `text` fields.

This subcommand is valuable in three scenarios:

- **CI/CD validation:** after model training, a pipeline step can run `flama model run` on a reference dataset and compare outputs to expected values.
- **Batch scoring:** for workloads where request-by-request API calls add unnecessary overhead, predictions can be computed in a single pass over a JSON file.
- **Debugging:** during development, a practitioner can test a model's behaviour on specific inputs without starting a server.

19.6.3 flama model stream

The `stream` subcommand performs streaming generation, printing tokens as they are produced. It supports both ML and LLM models, though streaming is most useful for LLMs where generation takes appreciable time:

```
$ echo "Write a haiku about Rust" \
  | flama model assistant.flm stream \
    --transport chat --channel all
```

The options mirror those of `run` (`--transport`, `--system`, `--param`, `--channel`), with the addition of `--buffer` which accumulates all output and writes it at once rather than printing incrementally.

This is particularly useful for interactive experimentation with LLMs during development, as it provides immediate visual feedback on generation quality and speed without requiring a running server.

19.7 flama upgrade

The `upgrade` command assists with codebase migration across major FLAMA versions. It rewrites import statements and renamed symbols across all Python files in the specified paths, applying automated transformations that cover the majority of breaking changes between versions.

```
# Preview changes as a unified diff (default)
$ flama upgrade src/ tests/

# Apply changes in place
$ flama upgrade --write src/ tests/
```

The command operates in two modes: `--diff` (default) previews the changes as a unified diff and exits with code 1 if any changes would be made, making it suitable for CI integration; `--write` applies the changes in place.

Option	Default	Type	Purpose
<code>-to</code>	(latest)	version	Target version
<code>-from</code>	(detect)	version	Source version to migrate from
<code>-diff/-write</code>	diff	flag	Preview or apply changes
<code>-select</code>	(all)	csv	Run only specified operations
<code>-skip</code>	(none)	csv	Skip specified operations

Symbols that have no automatic replacement are flagged with a `# flama-upgrade` marker and listed as manual follow-ups, ensuring that no breaking change passes silently.

20 Configuration and deployment

One application has to run in development, in staging, and in production, at different addresses, with different worker counts, different SSL material, and different model paths. Configuration is what absorbs that variation, and FLAMA splits it in two: what the application is (title, version, the models it serves) is kept apart from how it is served (host, port, workers). An application can be built from a Python import string, from a dictionary, or from a JSON file, and the three paths converge on the same objects. What follows covers those objects, the settings reader an application uses for its own configuration, and the deployment patterns the framework supports.

20.1 Configuration architecture

The configuration system consists of three dataclasses that separate application concerns from server concerns:

App encapsulates the application identity (title, version, description), feature configuration (schema route, docs route, debug mode), and the list of models to serve. The **App** class supports three construction paths: from a Python import string (`StrApp("mymodule:app")`), from a dictionary (`DictApp.from_dict(data)`), or from a live FLAMA instance (`FlamaApp(app)`).

Uvicorn wraps all Uvicorn server options into a single data structure: bind address, port, number of workers, SSL configuration, event loop selection, protocol implementations, reload settings, logging configuration, and timeouts. Every field has a sensible default and can be overridden from the CLI, from environment variables, or from the JSON configuration file.

Config combines an **App** and a **Uvicorn** instance and provides the `run()` method that starts the server. Config objects can be serialized to JSON (`config.dumps()`) and deserialized from JSON (`Config.loads(data)`) or from a file handle (`Config.load(fs)`).

20.2 Application settings

The dataclasses above describe how the CLI is told what to run. They are distinct from `flama.config`, which is what an application uses to read its *own* settings: credentials, feature flags, database coordinates, and anything else that varies between environments. A **Config** object is bound to an optional configuration file, in `ini`, `json`, `yaml`, or `toml` format, and resolves each key by looking first at the environment, then at that file, and finally at an explicit default; a key that is found nowhere and has no default raises `KeyError` rather than silently yielding `None`.

```
1 import dataclasses
2 from flama.config import Config, Secret
3
4 config = Config("config.yaml", format="yaml")
5
6 DEBUG = config("DEBUG", cast=bool)
7 API_KEY = config("API_KEY", cast=Secret)
8 DB_HOST = config("DATABASE.host")
9 TIMEOUT = config("TIMEOUT", default=30, cast=int)
10
11 @dataclasses.dataclass
12 class FeatureFlags:
13     enable_new_ui: bool
14     max_daily_limit: int
```

```

15
16 FEATURES = config("FEATURE_FLAGS", cast=FeatureFlags)

```

The `cast` argument accepts any type or callable, and two cases are worth singling out. A dotted key (`"DATABASE.host"`) walks into nested structures in the configuration file, so a hierarchical document can be read a leaf at a time. Casting to a dataclass parses the value as JSON when it arrives as a string, which is what an environment variable always is, and then constructs the dataclass from the matching keys, ignoring any others; this makes a structured setting expressible either as a nested block in the file or as a single JSON-valued environment variable, with the application reading it the same way in both cases.

`Secret` wraps a value so that printing it, logging it, or including it in a traceback shows `Secret('*****')` rather than the value itself, which remains available through `str()`. It is a guard against accidental disclosure in diagnostics, not an encryption mechanism.

20.3 Application generation

When used with the `serve` or `start` commands, the configuration system generates a temporary FLAMA application from a Jinja2 template. The template creates a FLAMA instance with the configured metadata and registers each model:

```

1 from flama import Flama
2
3 app = Flama(
4     debug=True,
5     openapi={
6         "info": {
7             "title": "ML Platform",
8             "version": "2.0.0",
9             "description": "Production model serving",
10        }
11    },
12    schema="/schema/",
13    docs="/docs/"
14 )
15
16 models = [{"url": "/sentiment", "path": "sentiment.flm",
17            "name": "sentiment"}]
18 for model in models:
19     app.models.add_model(
20         path=model["url"],
21         model=model["path"],
22         name=model["name"],
23     )

```

The generated module is written to a temporary file and passed to Uvicorn as an import path. This approach means that the `serve` and `start` commands produce fully standard FLAMA applications that can be inspected, debugged, and extended as needed.

20.4 Debug mode

When debug mode is enabled (`--app-debug` or `debug=True`), the framework activates the `ServerErrorMiddleware` in its interactive mode. Unhandled exceptions produce HTML error pages that include:

- The traceback, one entry per frame, each carrying the filename, the function, the line number, and ten lines of surrounding source.
- A marker on each frame recording whether it belongs to the application or to a vendored dependency, so the frames worth reading can be told from the ones that are only passing the exception along.
- The request that caused it: path, method, query parameters, headers, cookies, and the client's host and port.

The page is served only when the request's `Accept` header asks for HTML; anything else gets a plain `Internal Server Error`, which keeps API clients from receiving a page of source code. The exception is re-raised after the response is sent either way, so the server still logs it and a test client can still catch it.

Debug mode has no place in production. Source, file paths, and request headers are all exposed by it, and the framework does not currently warn about this at startup, so the responsibility for not shipping `debug=True` rests with the deployment configuration.

20.5 Deployment patterns

20.5.1 Direct server

The simplest deployment runs the FLAMA application directly with Uvicorn:

```
$ flama run myapp:app \  
  --server-host 0.0.0.0 \  
  --server-port 8000 \  
  --server-workers 4
```

This is suitable for internal services, development servers, and containerized deployments behind a load balancer.

20.5.2 Container deployment

A minimal Dockerfile for a FLAMA model server:

```
FROM python:3.12-slim  
WORKDIR /app  
COPY requirements.txt .  
RUN pip install --no-cache-dir -r requirements.txt  
COPY models/ models/  
COPY flama.json .  
EXPOSE 8000  
CMD ["flama", "start", "flama.json"]
```

The `flama start` command reads the configuration from `flama.json`, which specifies the models, routes, and server options. Environment variables can override any option at runtime, making the same container image deployable in different environments:

```
$ docker run -e HOST=0.0.0.0 -e PORT=8080 \
  -v /models:/app/models myimage:latest
```

20.5.3 Reverse proxy

For production deployments, FLAMA applications are typically placed behind a reverse proxy (Nginx, Traefik, or a cloud load balancer) that handles TLS termination, rate limiting, and static file serving. The `TrustedHostMiddleware` validates the `Host` header against a whitelist, and the `CORSMiddleware` manages cross-origin access policies.

21 Lifespan management

FLAMA implements the ASGI lifespan protocol, which provides a structured way to run initialization code at server startup and cleanup code at server shutdown. This replaces the ad-hoc signal handlers and module-level initialization patterns common in WSGI applications.

FLAMA provides three complementary mechanisms for managing application lifecycle:

Event handlers are individual async callables registered for the `startup` or `shutdown` events. They can be provided at construction time or registered dynamically with decorators.

Lifespan context manager is a callable that receives the application and returns an async context manager. The code before `yield` runs after all startup event handlers, and the code after `yield` runs before all shutdown event handlers.

Module hooks are `on_startup()` and `on_shutdown()` methods on `Module` subclasses. They are registered automatically when the module is added to the application.

Event handlers can be registered at construction time:

```
1 from flama import Flama
2
3 notifications = NotificationService()
4
5 async def connect_notifications():
6     await notifications.connect()
7
8 async def close_notifications():
9     await notifications.disconnect()
10
11 app = Flama(
12     events={
13         "startup": [connect_notifications],
14         "shutdown": [close_notifications],
15     }
16 )
```


Alternatively, the decorator syntax registers handlers after construction. An application using the `SQLAlchemyModule` typically uses a startup handler to create its schema, reaching the engine through the module rather than through any application-level state bag:

```
1 app = Flama(modules=[SQLAlchemyModule(DATABASE_URL)])
2
3 @app.on_event("startup")
4 async def create_schema():
5     async with app.sqlalchemy.engine.begin() as connection:
6         await connection.run_sync(metadata.create_all)
```

For applications that prefer the context-manager pattern, the `lifespan` parameter accepts a callable that returns an async context manager:

```
1 from contextlib import asynccontextmanager
2 from flama import Flama
3
4 @asynccontextmanager
5 async def lifespan(app):
6     # Runs after startup event handlers
7     model = load_heavy_model()
8     setattr(app, "model", model)
9     yield
10    # Runs before shutdown event handlers
11    delattr(app, "model")
12
13 app = Flama(lifespan=lifespan)
```

All startup event handlers run concurrently (via `run_task_group`), then the lifespan context manager is entered. On shutdown, the lifespan context manager exits first, then all shutdown event handlers run concurrently. FLAMA also propagates lifespan events to child mounted applications, ensuring that sub-applications initialise and tear down correctly.

Whatever a lifespan or a module sets up is reachable from the application object, which is itself injectable, so a component is all that is needed to hand it to request handlers as a typed dependency:

```
1 class DatabaseEngineComponent(Component):
2     def resolve(self, app: Flama) -> AsyncEngine:
3         return app.sqlalchemy.engine
```

FLAMA has no general-purpose state bag on the application; state belongs either to a module, which owns it behind its own name (Section 3.4), or to a component, which produces it on demand. Both are reached through the application instance, and both are typed.

22 Testing

FLAMA ships its own test client, `flama.client.Client`. It subclasses `httpx.AsyncClient`, so the request API is the familiar one, and wraps the application's lifespan: entering the context manager runs the startup handlers and module hooks, leaving it runs the shutdown ones. This matters more than convenience. A FLAMA application refuses to serve requests before its lifespan has run, so driving it through a bare `httpx.ASGITransport`, which does not implement the lifespan protocol, fails on the first request rather than exercising the application:

```

1 import pytest
2
3 from flama.client import Client
4 from myapp import app
5
6 @pytest.fixture(scope="function")
7 async def client():
8     async with Client(app=app) as client:
9         yield client
10
11 async def test_get_user(client):
12     response = await client.get("/users/1/")
13     assert response.status_code == 200
14     assert response.json()["name"] == "Ada Lovelace"
15
16 async def test_create_user(client):
17     response = await client.post(
18         "/users/",
19         json={"name": "Alan Turing",
20             "email": "alan@example.com"},
21     )
22     assert response.status_code == 201
23     assert response.json()["id"] is not None

```

Because resource routes are named, a test need not hard-code URLs: `app.resolve_url()` builds them from the resource and operation, so a change of mount point does not ripple through the suite.

```

1 async def test_list_users(client):
2     url = client.app.resolve_url("user:list").path
3     response = await client.get(str(url), params={"page_size": 100})
4     assert response.status_code == 200

```

22.1 Testing ML model endpoints

Model endpoints can be tested by mounting the application with a model file and sending prediction requests:

```

1 import pytest
2
3 from flama.client import Client
4
5 @pytest.fixture(scope="function")
6 async def ml_client():
7     async with Client(
8         models=[("classifier", "/classifier/",
9             "tests/fixtures/classifier.flm")]
10     ) as client:
11         yield client
12
13 async def test_model_inspect(ml_client):
14     response = await ml_client.get("/classifier/")
15     assert response.status_code == 200

```

```

16     assert response.json()["meta"]["framework"]["lib"] == "sklearn"
17
18     async def test_model_predict(ml_client):
19         response = await ml_client.model_request(
20             "classifier", "POST", "/predict/",
21             json={"input": [[5.1, 3.5, 1.4, 0.2]]},
22         )
23         assert response.status_code == 200
24         assert "output" in response.json()

```

The `models` argument is a shortcut that builds an application and registers each `(name, url, path)` triple through `add_model()`, which saves constructing one by hand when the models are all the test is about. `model_request()` then addresses a model by name, resolving the rest of the URL from where it was mounted.

22.2 Testing with Workers and repositories

Applications that use the Worker and Repository patterns (Section 7) can be tested with a real in-memory database or with mocked repositories:

```

1  import pytest
2
3  from flama import Flama
4  from flama.client import Client
5  from flama.sqlalchemy import SQLAlchemyModule, metadata
6
7  DATABASE_URL = "sqlite+aiosqlite:///memory:"
8
9  @pytest.fixture(scope="function")
10     async def client():
11         app = Flama(modules=[SQLAlchemyModule(DATABASE_URL)])
12         app.resources.add_resource("/users/", UserResource)
13
14         async with Client(app=app) as client:
15             engine = client.app.sqlalchemy.engine
16             async with engine.begin() as connection:
17                 await connection.run_sync(metadata.create_all)
18         yield client

```

The engine is taken from the `SQLAlchemyModule` once the client has entered the application's lifespan, which is what creates it; creating a second engine by hand would leave the tests talking to a different database from the application. An in-memory SQLite database gives each test run its own isolated schema without an external server.

For suites that must not pay for schema creation per test, the same structure works one level up: create the schema once in a session-scoped fixture, then wrap each test in a transaction that is rolled back on teardown. The module exposes `open_connection()`, `begin_transaction()`, `end_transaction()`, and `close_connection()` for exactly that, so a test can share the application's connection and still leave no trace behind.

Part V

Ecosystem and Outlook

23 Related work

Python web frameworks and ML serving systems have evolved along tracks that barely touch. Web frameworks took up asynchronous execution and type-driven validation and largely stopped there, while model serving platforms grew up separately to move trained models into production, bringing packaging formats, batching, and monitoring but almost none of the HTTP machinery. A third track opened more recently with the LLM inference engines built for generative workloads specifically. Each is worth taking in turn, because what FLAMA is depends on what these are not.

23.1 Web frameworks

Flask ([Ronacher, 2010](#)) has been the default answer for small and medium Python web applications since 2010, and its minimalism is deliberate: URL routing, template rendering, and a request/response abstraction are provided, and everything else, validation and serialization and authentication and database access alike, is left to extensions. That works until concurrency enters the picture. Being synchronous and WSGI-based, Flask has no answer of its own for I/O-bound or CPU-bound work beyond an external queue such as Celery or RQ, which is a second piece of infrastructure to run and monitor. Nothing in it is ML-specific.

Django ([Django Software Foundation, 2005](#)) takes the opposite position and supplies almost everything: an ORM, an admin interface, authentication, templates, with the Django REST Framework ([Christie, 2011](#)) adding serialization, viewsets, and generated API documentation on top. The cost is weight and a synchronous ORM whose async support remains partial, which tells against it for lightweight API services and real-time workloads. On models it is as silent as Flask.

Starlette ([Christie, 2018](#)) is a lightweight ASGI toolkit supplying routing, request and response classes, middleware, WebSocket support, and a test client. It is a toolkit rather than a framework: it does not include validation, dependency injection, or API documentation. FLAMA's 1.x series was built on top of Starlette; the 2.0 rewrite replaced it with a native HTTP, routing, and middleware core, part of it Rust-accelerated (Section 3.1), removing the dependency.

23.2 ML serving platforms

TensorFlow Serving ([Google, 2016](#)) is very good at one thing. Given a SavedModel artifact it will version it, batch requests against it, and expose it over gRPC or REST at high throughput. What it will not do is anything else: the models have to be TensorFlow's, the endpoints are the ones it generates, and there is no route by which an inference result might be joined to a database row before being returned.

TorchServe ([PyTorch Team, 2020](#)) occupies the same position for PyTorch, with model archiving, worker management, metrics, and A/B testing. Custom handler classes provide some room to

manoeuvre, but the shape of the deployment is fixed: a standalone server that application logic has to be arranged around rather than composed with.

Serving is only a part of what **MLflow** ([Databricks, 2018](#)) does, most of it being experiment tracking and a model registry. A model can be put behind a REST API from there, though the component that does it is a thin wrapper, and none of the routing, validation, authentication, or database integration that a web framework would bring comes with it.

Closest of all in ambition is **BentoML** ([BentoML Team, 2019](#)), which has a packaging format of its own (Bento), batching and concurrency in its serving layer, and support across ML frameworks. The divergence is in what that serving layer is for. It is built for model inference specifically, so pluggable schema validation, dependency injection, database-backed resources, domain-driven design, and JWT authentication are outside its remit, and an application needing both inference and ordinary API endpoints ends up running BentoML alongside a web framework rather than instead of one.

23.3 LLM inference engines

vLLM ([Kwon et al., 2023](#)) is a high-throughput inference engine for large language models. It introduces PagedAttention for efficient KV-cache management, continuous batching for maximizing GPU utilization, and tensor parallelism for multi-GPU deployments. An OpenAI-compatible API server comes with it, which is often mistaken for the thing being a web framework. It is not one. Schema validation, dependency injection, resource generation, authentication, database access, and any endpoint that is not an LLM endpoint all fall outside it, so an application wanting inference and an ordinary API runs vLLM beside itself as a separate service. FLAMA uses vLLM as a backend for exactly this reason: the engine is worth having, and the framework is a different problem.

Ollama ([Ollama, 2023](#)) aims lower and hits its target, handling model download, quantization, and local serving behind a REST API of its own design. The composition limits are the same as vLLM's, and for the same structural reason.

LiteLLM ([BerriAI, 2023](#)) is neither engine nor framework but a translator, converting between the OpenAI, Anthropic, Ollama, and other formats so that an application can change provider without changing code. It runs no models itself. Where FLAMA renders several dialects from one in-process model, LiteLLM forwards to a provider that does the rendering, which is a genuine alternative to multi-dialect serving and costs a network hop per request.

23.4 Positioning Flama

Three categories come out of the survey. There are web frameworks that can be extended toward model serving (Flask, Django), predictive model servers that can be extended a little way toward the web (TensorFlow Serving, TorchServe, BentoML), and LLM engines presenting a single protocol (vLLM, Ollama). What they share is the direction of the gap: whichever one is picked, the others' capabilities arrive later as external tools, glue code, or an operational workaround.

FLAMA sits where none of them do, as a web framework whose model serving is native rather than bolted on, for predictive and generative models alike. The claim is narrower than it sounds and rests on a single implementation fact: the routing, validation, dependency injection, middleware, and authentication that a conventional endpoint goes through are not merely similar to what an inference endpoint goes through, they are the same objects. A prediction route is a route. It carries the same tags the authentication middleware reads, and the paginator does not know or care whether the collection it wraps came out of a table or a model.

The multi-dialect layer follows from the same arrangement. One deployment answers the OpenAI SDK, the Anthropic SDK, and the Ollama CLI at once, with no protocol-specific application code and no proxy in between, because the dialects are renderers over one canonical event stream rather than separate servers.

24 Feature comparison

Table 1 summarizes the feature sets of FLAMA and six representative frameworks across three dimensions: web API capabilities, predictive ML serving, and LLM inference capabilities. Features are evaluated based on built-in support only; third-party extensions and custom code are not counted.

Feature	Flama	Flask	TF Serving	TorchServe	BentoML	vLLM	Ollama
ASGI / async-native	✓	–	–	–	–	–	–
WebSocket & streaming	✓	–	–	–	–	–	–
SSE / NDJSON responses	✓	–	–	–	–	✓	✓
Type-driven validation	✓	–	–	–	–	–	–
Pluggable schema libs	✓	–	–	–	–	–	–
Component-based DI	✓	–	–	–	–	–	–
CRUD resource gen.	✓	–	–	–	–	–	–
DDD patterns	✓	–	–	–	–	–	–
JWT authentication	✓	–	–	–	–	–	–
OpenAPI generation	✓	–	–	–	–	–	–
Multi-framework ML	✓	–	–	–	✓	–	–
Model packaging format	✓	–	–	–	✓	–	–
Codeless model serving	✓	–	✓	✓	✓	✓	✓
Multi-backend LLM	✓	–	–	–	–	–	–
Multi-dialect serving	✓	–	–	–	–	–	–
MCP support	✓	–	–	–	–	–	–
Model acquisition CLI	✓	–	–	–	–	–	✓
Codebase migration	✓	–	–	–	–	–	–
Unified API + ML + LLM	✓	–	–	–	–	–	–

Table 1 Feature comparison across frameworks, serving platforms, and LLM inference engines. A checkmark indicates built-in support without third-party extensions. FLAMA is the only system that provides full web API capabilities, native multi-framework ML model serving, multi-backend LLM inference with simultaneous multi-dialect exposure, and MCP support within a single unified architecture.

Several observations emerge from this comparison:

Read down the columns rather than across the rows and the pattern is a clean split. Everything to the right of FLAMA is strong in one band of the table and empty in the others, and nothing in the survey is strong in two. That is the observation the table exists to make.

Three qualifications are worth attaching to it. The first is that a checkmark records presence, not quality: vLLM’s PagedAttention is a better piece of inference engineering than anything FLAMA contributes, which is why FLAMA runs on top of it rather than against it, and the row marked *multi-backend LLM* should be read as coverage rather than as a claim about throughput. The second is that several of these systems are not trying to fill the other bands, so an empty cell is often a scope decision rather than an omission; TensorFlow Serving is not a worse web framework than

FLAMA, it is not one. The third is that the table rewards breadth, and breadth is only a virtue for an application that actually needs it. A service doing nothing but batch inference over one TensorFlow model is better served by TensorFlow Serving, and the argument made here does not say otherwise.

What the comparison does support is narrower: for an application that needs conventional endpoints and predictive inference and generative serving at once, the alternatives require composing two or three systems, and FLAMA requires one.

25 Conclusion and future work

This paper has presented FLAMA, an open-source Python framework that unifies web API development, predictive machine-learning model serving, and large-language-model inference under a single programming model. The framework rests on seven design principles: async-first execution on the ASGI standard, type annotations as the source of truth for validation and documentation, pluggable schema libraries for data validation, component-based dependency injection, convention over configuration for resource generation, native ML model serving for the four major predictive frameworks, and protocol-agnostic LLM serving with multi-backend inference and multi-dialect wire format support.

The FLM binary format provides a portable, compressed, metadata-rich container for both predictive models and LLM checkpoints, with two protocol versions addressing the distinct storage requirements of each family. The three-level predictive integration system (components, `add_model()`, and model resources) and the multi-dialect LLM serving architecture (backends, transport, codec, dialects) allow developers to choose the degree of customization that matches their deployment requirements. The MCP module, now implementing the current stateless protocol revision with Tasks, Elicitation, and Apps extensions, connects FLAMA applications to the broader AI tool ecosystem.

FLAMA is in production use for conventional REST APIs, real-time inference services, and multi-dialect LLM endpoints. Underneath sits the Rust core, seven modules covering route resolution, path and host matching, JSON encoding, compression, multipart parsing, cookie handling, and HTTP utilities, and the effect of having it is visible in the numbers rather than only in the design: growing a route table from ten entries to two hundred moves the cost of a request by roughly 4%, and ten layers of middleware add about 0.6% over none at all, both measured under Callgrind in continuous integration. The point of compiling those paths was to stop the framework's own overhead from being the thing an application has to plan around.

Several directions for future work are planned:

- **Model versioning and A/B testing.** Built-in support for serving multiple model versions concurrently and routing traffic between them based on configurable policies (percentage splits, header-based routing, canary deployments).
- **Adaptive batching.** Automatic request batching for GPU-bound inference to maximize throughput without requiring user configuration. Incoming prediction requests would be accumulated into micro-batches based on latency budgets and batch size targets.
- **GraphQL support.** Extending the schema and routing system to support GraphQL endpoints alongside REST, with the same type-driven validation and dependency injection infrastructure.
- **Observability integration.** Structured logging, distributed tracing (OpenTelemetry), and metrics export as built-in middleware, providing production-grade observability without third-party

instrumentation.

- **Agentic workflows.** Native support for multi-step agent orchestration, combining LLM generation with tool use in structured loops, with built-in state management and conversation persistence.

FLAMA is released under the Apache 2.0 licence. The source code is available at <https://github.com/vortico/flama>, with documentation at <https://flama.dev> and package distribution via PyPI (`pip install flama`).

References

- Martín Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, et al. TensorFlow: A system for large-scale machine learning. In *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*, pages 265–283. USENIX Association, 2016. URL <https://www.usenix.org/conference/osdi16/technical-sessions/presentation/abadi>.
- Anthropic. Anthropic messages API. <https://docs.anthropic.com/en/api>, 2024. Accessed: 2026-06-15.
- Anthropic. Model context protocol specification, revision 2026-07-28. <https://modelcontextprotocol.io/specification/2026-07-28>, 2026. Accessed: 2026-08-09.
- Michael Bayer. SQLAlchemy: The Python SQL toolkit and object relational mapper. <https://www.sqlalchemy.org>, 2006. Accessed: 2026-04-13.
- BentoML Team. BentoML: The unified model serving framework. <https://www.bentoml.com>, 2019. Accessed: 2026-04-13.
- BerriAI. LiteLLM: Call all LLM APIs using the OpenAI format. <https://github.com/BerriAI/litellm>, 2023. Accessed: 2026-06-15.
- Tom Christie. Django REST Framework. <https://www.django-rest-framework.org>, 2011. Accessed: 2026-04-13.
- Tom Christie. Uvicorn: An ASGI web server for Python. <https://uvicorn.dev>, 2017. Accessed: 2026-08-19.
- Tom Christie. Starlette: The little ASGI framework that shines. <https://www.starlette.io>, 2018. Accessed: 2026-04-13.
- Tom Christie. Typesystem: Data validation and form rendering for Python. <https://www.encode.io/typesystem>, 2019. Accessed: 2026-04-13.
- Yann Collet and Murray Kucherawy. Zstandard compression and the `application/zstd` media type. RFC 8878, 2021.
- Samuel Colvin. Pydantic: Data validation using Python type annotations. <https://docs.pydantic.dev>, 2017. Accessed: 2026-04-13.
- Databricks. MLflow: An open source platform for the machine learning lifecycle. <https://mlflow.org>, 2018. Accessed: 2026-04-13.
- Django Software Foundation. Django: The web framework for perfectionists with deadlines. <https://www.djangoproject.com>, 2005. Accessed: 2026-04-13.
- Phillip J. Eby. PEP 333 – Python web server gateway interface v1.0. <https://peps.python.org/pep-0333/>, 2003. Python Enhancement Proposal, Informational, created 7 December 2003. Accessed: 2026-08-19.
- Eric Evans. *Domain-Driven Design: Tackling Complexity in the Heart of Software*. Addison-Wesley, 2003. ISBN 978-0-321-12521-7.
- I. Fette and A. Melnikov. The WebSocket protocol. RFC 6455, 2011. Standards Track, December 2011.

- Martin Fowler. *Patterns of Enterprise Application Architecture*. Addison-Wesley, 2002.
- Martin Fowler. Inversion of control containers and the dependency injection pattern. <https://martinfowler.com/articles/injection.html>, 2004. Published 23 January 2004. Accessed: 2026-08-19.
- Andrew Godwin and ASGI contributors. ASGI (asynchronous server gateway interface) specification. <https://asgi.readthedocs.io>, 2016. Accessed: 2026-04-13.
- Ryan Gonzalez, Philip House, Ivan Levkivskyi, Lisa Roach, and Guido van Rossum. PEP 526 – syntax for variable annotations. <https://peps.python.org/pep-0526/>, 2016. Python Enhancement Proposal, Standards Track, created 9 August 2016. Accessed: 2026-08-19.
- Google. TensorFlow Serving: Flexible, high-performance serving system for machine learning models. <https://www.tensorflow.org/tfx/guide/serving>, 2016. Accessed: 2026-04-13.
- Awni Hannun, Jagrit Digani, Angelos Katharopoulos, and Ronan Collobert. MLX: Efficient and flexible machine learning on apple silicon. <https://github.com/ml-explore/mlx>, 2023. Accessed: 2026-06-15.
- M. Jones, J. Bradley, and N. Sakimura. JSON Web Signature (JWS). RFC 7515, 2015a. Standards Track, May 2015.
- M. Jones, J. Bradley, and N. Sakimura. JSON Web Token (JWT). RFC 7519, 2015b.
- JSON-RPC Working Group. JSON-RPC 2.0 specification. <https://www.jsonrpc.org/specification>, 2013. Origin date 26 March 2010, updated 4 January 2013. Accessed: 2026-08-19.
- Keras Team. Save, serialize, and export models: The .keras format. https://keras.io/guides/serialization_and_saving/, 2023. Accessed: 2026-08-19.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP '23)*, pages 611–626. ACM, 2023. doi: 10.1145/3600006.3613165.
- Łukasz Langa. PEP 585 – type hinting generics in standard collections. <https://peps.python.org/pep-0585/>, 2019. Python Enhancement Proposal, Standards Track, created 3 March 2019. Accessed: 2026-08-19.
- Steven Loria. marshmallow: simplified object serialization. <https://marshmallow.readthedocs.io>, 2013. Accessed: 2026-04-13.
- NDJSON Working Group. NDJSON: Newline delimited JSON. <https://github.com/ndjson/ndjson-spec>, 2014. Specification version 1.0.0, last updated 19 October 2014. Accessed: 2026-08-19.
- Ollama. Ollama: Get up and running with large language models locally. <https://github.com/ollama/ollama>, 2023. Accessed: 2026-06-15.
- OpenAI. OpenAI API reference. <https://platform.openai.com/docs/api-reference>, 2023. Accessed: 2026-06-15.
- OpenAPI Initiative. OpenAPI specification, version 3.2.0. <https://spec.openapis.org/oas/v3.2.0>, 2025. Released 19 September 2025. Accessed: 2026-08-19.

- Andrei Paleyes, Raoul-Gabriel Urma, and Neil D. Lawrence. Challenges in deploying machine learning: A survey of case studies. *ACM Computing Surveys*, 55(6):114:1–114:29, 2022. doi: 10.1145/3533378.
- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems*, volume 32, pages 8024–8035. Curran Associates, Inc., 2019.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- Philippe Prados and Maggie Moss. PEP 604 – allow writing union types as `X | Y`. <https://peps.python.org/pep-0604/>, 2019. Python Enhancement Proposal, Standards Track, created 28 August 2019. Accessed: 2026-08-19.
- PyO3 Project. Maturin: Build and publish crates with pyO3, cffi and uniffi bindings as well as Rust binaries as Python packages. <https://github.com/PyO3/maturin>, 2019. Accessed: 2026-06-15.
- PyTorch Team. TorchServe: A flexible and easy to use tool for serving PyTorch models. <https://pytorch.org/serve>, 2020. Accessed: 2026-04-13.
- PyTorch Team. `torch.export`: Ahead-of-time export of PyTorch models to `ExportedProgram`. <https://docs.pytorch.org/docs/stable/export.html>, 2024. Accessed: 2026-08-19.
- Armin Ronacher. Flask: A micro web framework for Python. <https://flask.palletsprojects.com>, 2010. Accessed: 2026-04-13.
- Armin Ronacher. Click: A composable command line interface toolkit for Python. <https://click.palletsprojects.com>, 2014. Accessed: 2026-04-13.
- D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-François Crespo, and Dan Dennison. Hidden technical debt in machine learning systems. In *Advances in Neural Information Processing Systems*, volume 28, pages 2503–2511, 2015.
- Guido van Rossum, Jukka Lehtosalo, and Łukasz Langa. PEP 484 – type hints. <https://peps.python.org/pep-0484/>, 2014. Python Enhancement Proposal, Standards Track, created 29 September 2014. Accessed: 2026-08-19.
- Till Varoquaux and Konstantin Kashin. PEP 593 – flexible function and variable annotations. <https://peps.python.org/pep-0593/>, 2019. Python Enhancement Proposal, Standards Track, created 26 April 2019. Accessed: 2026-08-19.
- W3C. Trace context, level 1. <https://www.w3.org/TR/trace-context/>, 2021. W3C Recommendation, 23 November 2021. Accessed: 2026-08-19.
- WHATWG. Server-sent events. <https://html.spec.whatwg.org/multipage/server-sent-events.html>, 2026. HTML Living Standard, continuously updated. Accessed: 2026-06-15.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, et al. Transformers: State-of-the-art

natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45. Association for Computational Linguistics, 2020. doi: 10.18653/v1/2020.emnlp-demos.6.