

Daedalus-150M: A Convolution–Attention Hybrid Designed for CPU Inference

Christos Koutsiaris
Development Expert

Abstract—Small language models are usually built like large ones and then squeezed onto a CPU afterwards. We did the opposite: we fixed the target first — one user, one token at a time, 4-bit weights, ordinary CPU — and chose the architecture to suit it. The result keeps full attention in only 6 of its 18 blocks. The other 12 use short convolutions whose memory is two timesteps wide no matter how long the conversation gets, so two thirds of the network never re-reads a growing cache.

Trained from scratch on 59.9 B tokens, the model scores 47.31 on a five-task benchmark against a bar of 42.20 that was fixed before training began. It beats GPT-2 124M, Pythia-160M, OPT-125M and GPT-neo-125M, all trained on three to six times more data, and exceeds MobileLLM-125M’s published score despite that model seeing a trillion tokens. Validation bits-per-byte is 0.8685.

To check the architecture rather than the training recipe, we trained a conventional all-attention model of the same size on the same data, and wrote down the winning condition before scoring either. The hybrid won the chosen quality metric by 0.81 %, matched it on downstream tasks, produced a 6.3 % smaller 4-bit file, and decoded $1.76\times$ faster at 2048 tokens of context — $2.08\times$ against an external model of similar size. In every measurement the speed advantage is near zero at an empty context and grows with length, which is what the mechanism predicts and what a merely leaner model would not show. A simple bandwidth calculation predicts only $1.17\times$, so memory volume alone does not explain the gap.

We also report what did not work: an unmitigated 4-bit quality cost, roughly half the convolution channels ending up inert and impossible to remove, and a vocabulary larger than this model size warrants.

Index Terms—language models, CPU inference, efficient architectures, hybrid models, quantisation, memory bandwidth, edge deployment

I. INTRODUCTION

A language model serving a single user on a CPU is not a smaller version of a model serving many users on a GPU. Three properties of the regime differ, and together they change what an efficient architecture looks like.

Batch size is one. There is no batching over which to amortise weight loading. Every decoded token requires streaming the model’s weights through the memory hierarchy, so throughput is governed by bytes read per token rather than by arithmetic throughput.

Memory bandwidth binds before compute. A contemporary CPU can issue far more arithmetic than it can feed with operands. This inverts the usual optimisation target: an architecture performing more arithmetic over fewer bytes runs faster.

The key–value cache is a growing tax. In an all-attention decoder, generating each token requires re-reading the keys and values of every preceding token in every layer. That cost is linear in context length and is paid per token. Under high-bandwidth, large-batch GPU serving it is tolerable; at batch size one on a CPU it dominates long-context decoding.

The last property is the design lever this work pulls. If most layers carry a *constant-size* state rather than a growing one, decode cost becomes far flatter in context length, and the advantage widens precisely where users perceive latency: long conversations and long documents.

A. Contributions

- 1) A concrete sub-200M architecture whose layer composition is chosen from a CPU memory-traffic argument rather than adapted from GPU practice.
- 2) A pre-registered, parameter-matched ablation isolating the effect of that composition, with the decision rule fixed before either arm was scored.
- 3) Decode measurements at three context depths against both an internal twin and an external peer, showing the advantage grows with depth in both.
- 4) A first-order cost model that under-predicts the measured advantage, localising the residual to latency rather than bandwidth.
- 5) A negative result closing structural pruning of dead convolution channels, and a full account of five deviations forced on the training run.

B. Success criteria fixed in advance

To prevent post-hoc rationalisation, the evaluation bar was fixed by operator decision before any headline number existed. Quality is the five-task mean over HellaSwag, ARC-Easy, PIQA, OpenBookQA and WinoGrande, scored under lm-evaluation-harness conventions on a single harness so that peers and this model are measured identically. Table I lists the comparison set.

The bar is **42.2**, the strongest of the four beatable peers. Clearing it means matching or exceeding 300 B-token-class models on roughly seven times less data. The second half of the bar—CPU decode speed—is an architectural property and does not depend on training outcome.

II. RELATED WORK

Small language models. A line of work has shown that sub-1 B models improve far past classical compute-optimal token

TABLE I
PEER SET AND THE PRE-REGISTERED BAR. SCORES ARE MEASURED ON
THIS PROJECT’S HARNESS, NOT QUOTED FROM PUBLICATIONS.

Model	Tokens	5-task	Target
Pythia-160M	300 B	41.0	beat
GPT-neo-125M	300 B	41.9	beat
OPT-125M	180 B	42.1	beat
GPT-2 124M	—	42.2	the bar
MobileLLM-125M	1 T	gated	stretch
Peer-135M	2 T	51.2	concede quality

budgets when trained on carefully filtered data. Pythia [6] established reproducible small baselines; MobileLLM [7] showed that depth-over-width and embedding sharing matter disproportionately below 350 M parameters, a finding this design adopts in its narrow feed-forward and tied embeddings.

Recurrent and hybrid sequence models. Structured state-space models [8] and gated linear-recurrent architectures [9] replace attention with mechanisms whose inference state is constant in sequence length. Purely recurrent models, however, lose the precise associative recall that attention provides. Hybrids that retain a minority of attention layers—Griffin [10] being a prominent example—recover most of that capability while keeping the majority of layers cache-free. Daedalus sits in this family, with two distinguishing choices: the recurrent operator is a very short depthwise convolution rather than a selective-scan state-space layer, chosen because it maps onto existing CPU inference kernels without new operators; and the attention-to-recurrence ratio is fixed by a memory-traffic argument specific to batch-size-one CPU decoding.

Efficient attention. Multi-query and grouped-query attention [3] shrink the cache by sharing key-value heads across query heads. This is complementary rather than alternative: Daedalus applies GQA to the six attention layers it retains, compounding the two reductions.

Quantisation. Post-training quantisation to 4 bits is standard for CPU deployment. We target the reference runtime’s [5] Q4_0 format specifically because its dot-product kernels are the best optimised on the target hardware, rather than a format with better theoretical error at lower throughput.

What is new here. The contribution is not a new operator. It is the observation that in the batch-size-one CPU regime the *ratio* of cache-free to attention layers is the dominant deployable design variable, together with a controlled experiment that isolates its effect at matched parameter count.

III. ARCHITECTURE

Daedalus-150M has 160.49 M parameters in 18 blocks at $d_{\text{model}} = 768$, feed-forward inner dimension 2048, vocabulary 49,152 and context 2048. Six blocks use full attention and twelve use short convolutions, interleaved as

C C C C A C C A C A C A C A C C A C

where A denotes attention (indices 4, 7, 9, 11, 13, 16) and C denotes convolution. Fig. 1 shows the stack and the internals of both block types.

A. The short-convolution block

Each convolution block computes

$$B, C, x = \text{in_proj}(u) \quad (\text{split three ways}) \quad (1)$$

$$y = \text{depthwise_conv1d}(B \odot x) \quad (2)$$

$$\text{out} = \text{out_proj}(C \odot y) \quad (3)$$

with kernel length $L = 3$ and groups equal to the channel count, so the convolution is depthwise and channels evolve independently. Its recurrent state is exactly $L - 1 = 2$ timesteps wide regardless of context length. The gating terms B and C supply the input-dependent behaviour a fixed kernel alone lacks.

This is the crux of the design: *decoding through a convolution block costs the same at token 2000 as at token 2, whereas decoding through an attention block does not.*

B. Why six attention layers

Pure convolution forfeits precise long-range retrieval, the operation attention performs uniquely well. The design retains six full-attention blocks, spread through the depth rather than clustered, so retrieval capacity is distributed across levels of representation.

The ratio is the whole trade. Against a 24-layer all-attention twin, the hybrid maintains a key-value cache in six layers rather than twenty-four.

C. Supporting choices

Grouped-query attention [3] with four key-value heads for twelve query heads reduces the cache on the six attention layers by a further factor of three. GQA is usually motivated by GPU memory capacity; here the motivation is bytes per token.

Tied embeddings. At $49,152 \times 768$ the embedding matrix holds 37.7 M parameters, 23 % of the model. Tying input and output projections removes an entire copy from both the artefact and memory traffic.

Feed-forward inner dimension 2048 ($2.67 \times d_{\text{model}}$ rather than the conventional $4\times$) shifts parameters into depth and attention and away from the widest, most bandwidth-hungry tensors.

D. Rejected alternatives

Mixture-of-experts routing was rejected at design time: no sub-1 B precedent exists at this scale, and the reference runtime’s mixture path hardcodes a specific vendor’s gating, so the model would not export. Positional-encoding-free attention was rejected because it breaks the export format. Distillation from a larger teacher was cancelled on budget grounds—storing teacher logits required approximately 288 GB against a 250 GB disk.




Fig. 1. Daedalus-150M. (a) The 18-block stack: 12 short-convolution blocks (light) interleaved with 6 attention blocks (dark) at indices 4, 7, 9, 11, 13 and 16, between a tied embedding and language-model head. (b) A convolution block carries a two-timestep state whose size is independent of context, so its decode cost is constant in t . (c) An attention block must re-read the whole key-value cache per generated token, so its cost is linear in t ; grouped-query attention (4 KV heads for 12 query heads) reduces the constant but not the growth. Every block also carries a feed-forward sub-layer; it is drawn once in (c) for space. The design question this paper answers empirically is what ratio of (b) to (c) is right when the deployment target is batch-size-one CPU decoding.

IV. A COST MODEL FOR CPU DECODING

Before presenting measurements we state what the architecture *should* buy, so that the measurement can disagree with it.

At batch size one, generating one token requires reading every weight once, plus the key-value cache of every attention layer over the current context. Let W be the quantised weight bytes, L_A the number of attention layers, h_{kv} the key-value heads, d_h the head dimension, b the bytes per cache element and t the context depth. Bytes read per generated token are

$$M(t) = W + \underbrace{2L_A h_{kv} d_h b}_{\kappa} t \quad (4)$$

where κ is the cache bytes per context token.

For the hybrid, $L_A = 6$, $h_{kv} = 4$, $d_h = 64$, $b = 2$, giving $\kappa_{\text{hyb}} = 6144\text{B}$. For the dense twin ($L_A = 24$, $h_{kv} = 2$, $d_h = 64$), $\kappa_{\text{dense}} = 12,288\text{B}$ — **exactly twice**, despite the twin’s narrower per-layer cache, because it has four times as many attention layers. Table II evaluates Eq. (4) against measurement.

The model is directionally right and quantitatively wrong, and the discrepancy widens with depth. Pure byte

TABLE II
FIRST-ORDER BANDWIDTH MODEL VERSUS MEASUREMENT. THE MODEL CAPTURES THE DIRECTION AND THE GROWTH BUT SUBSTANTIALLY UNDER-PREDICTS THE MAGNITUDE.

Depth	M hybrid	M dense	Predicted	Measured
0	100.2 MB	106.6 MB	1.06×	1.20×
512	103.3 MB	112.8 MB	1.09×	1.45×
2048	112.8 MB	131.7 MB	1.17×	1.76×

accounting predicts a 17% advantage at depth 2048; the measured advantage is 76%. We therefore reject the hypothesis that bandwidth alone explains the result and attribute the residual to two effects that Eq. (4) does not model. First, attention traverses the cache with a dependent softmax reduction, which is latency-bound rather than bandwidth-bound and scales poorly when the working set exceeds last-level cache; a depthwise convolution instead streams a two-element state with perfect locality. Second, the twin executes 24 layers to the hybrid’s 18, so per-layer fixed costs are paid a third more often.

The practical consequence is that the architecture’s benefit is *larger* than a naive byte count suggests, and that improving

TABLE III
CORPUS COMPOSITION (BLUEPRINT SHARES).

Source	Share	Role
FineWeb-Edu	37.5 %	educational web, backbone
DCLM-baseline	22.5 %	broad filtered web
Stack-Edu (Python)	9.0 %	code
FinePDFs-Edu	8.0 %	long-form document prose
FinePhrase	7.0 %	phrase-level diversity
Cosmopedia-v2	5.0 %	synthetic textbook style
FineMath-3+	3.0 %	mathematics
InfiWebMath-3+	3.0 %	web mathematics
FineWiki-en	3.0 %	encyclopedic reference
Everyday-conversations	2.0 %	dialogue register

attention kernels would narrow but not close the gap—the cache the hybrid does not keep cannot be optimised.

V. CORPUS AND DATA PIPELINE

Training data is a ten-source English mixture totalling **16.93 B unique tokens**, weighted toward educational and reasoning-dense text (Table III). At this scale data quality dominates quantity, and the evaluation suite is knowledge- and commonsense-heavy.

A. Bounded repetition

A 59.9 B-token budget over a 16.93 B-token corpus implies roughly 3.5 epochs. Rather than allow any source to be repeated arbitrarily to meet its nominal share, each source is capped at four epochs, following the finding that repetition up to approximately four epochs costs little relative to unique data [2]. Mass freed by a capped source is redistributed by water-filling across sources retaining headroom.

The cap is load-bearing. Without it, a source such as *everyday-conversations* (approximately 400 k tokens) would be repeated thousands of times to fill a 2 % share, which is not what a 2 % share is meant to express.

B. Packing and held-out data

Documents are tokenised and packed into fixed-size shards of contiguous token ids. Each source reserves whole shard *files*—never token slices—for a ≈ 2 % holdout, so no training window can straddle the train/holdout boundary. Because shard files are reserved whole, the realised holdout fraction varies per source: a source whose trailing shard is small relative to the target must reserve the preceding shard too, and one source’s real carve is 9.16 % rather than 2 %. Uneven carves perturb the training mixture, so the effect is measured rather than assumed.

Validation bits-per-byte is weighted by the probabilities the sampler actually draws with, not by holdout token counts. Correcting this distinction changed the reported metric by 9.7 %: weighting by holdout size systematically over-weighted the hardest sources, because those sources happened to have larger trailing shards.

TABLE IV
TRAINING CONFIGURATION.

Budget	59.9 B tokens, 124,476 steps
Optimisers	Muon (122.68 M params) + AdamW (37.81 M)
Learning rates	Muon 0.02, AdamW 3×10^{-4}
Schedule	WSD, 300-step warmup, linear decay to zero over the final 45 %
Batch ramp	128 k $\rightarrow$ 512 k tokens/step (first 10 %)
Sequence ramp	1024 $\rightarrow$ 2048 (first 10 %)
Regularisation	z -loss 10^{-4} , gradient clip 1.0
Precision	bf16
Hardware	1 $\times$ RTX 5090 (32 GB)

VI. TRAINING

Table IV summarises the configuration. Parameters are split by tensor shape: two-dimensional weight matrices are optimised by Muon [4], and embeddings, norms and biases by AdamW—122.68 M parameters over 102 tensors and 37.81 M over 62 respectively.

A. Decaying the learning rate to zero

The schedule is warmup–stable–decay with *linear decay to zero*, not cosine decay to a floor, following evidence that full decay to zero is substantially more sample-efficient [1].

A practical consequence deserves emphasis for anyone reading the loss curve: **training loss is expected to plateau during the stable phase**. Between approximately step 20,000 and step 68,461 the learning rate is held at its peak and loss moves very little. This is the schedule operating as designed, not a stall; nearly all remaining quality is purchased during the decay phase. In this run the plateau was mistaken for a stall by an observer, which is precisely why it is documented here.

A secondary benefit is that the checkpoint at the onset of decay is a reusable branch point: it can be trained further on additional or different data and re-decayed, whereas a model already annealed to $\text{lr} \approx 0$ requires learning-rate re-warmup.

VII. CENTRAL EXPERIMENT: HYBRID VERSUS DENSE

The architectural claim was tested directly, with the decision rule fixed before either arm was scored. Two models, parameter-matched to within 0.5 % (160.49 M hybrid against 161.25 M dense, the twin being 24 all-attention layers at $d = 640$ with feed-forward 2304), were trained on identical data and schedule for 5 B tokens each and fully decayed. The pre-registered metric was validation bits-per-byte over a 645 M-token held-out set with a 0.5 % margin floor fixed in advance. A dense win beyond that floor was a live possible outcome that would have changed the main run’s architecture.

A. Quality

The hybrid wins the pre-registered metric by 0.81 %, clearing the 0.5 % floor. On the five-task mean the dense twin is nominally ahead by 0.14 points, approximately 0.24σ against the $\approx 0.58\sigma$ these suites carry. Per task the two swap places in both directions—the hybrid leads on HellaSwag and

TABLE V
ABLATION QUALITY. THE TWO METRICS DISAGREE IN DIRECTION.

	val_bpb ↓	5-task
Daedalus-150M (hybrid)	0.910398	44.68
dense-150m (twin)	0.917774	44.82

TABLE VI
PER-SOURCE VALIDATION BITS-PER-BYTE (DENSE ARM, 5 B TOKENS).

Source	val_bpb	Weight
Stack-Edu (Python)	0.5811	9.2 %
Cosmopedia-v2	0.6109	5.1 %
FineWiki-en	0.7625	3.1 %
FinePDFs-Edu	0.9241	8.2 %
InfiWebMath-3+	0.9361	3.1 %
FinePhrase	0.9381	7.1 %
FineWeb-Edu	0.9451	38.3 %
FineMath-3+	0.9662	3.1 %
DCLM-baseline	1.0783	23.0 %

OpenBookQA, the twin on WinoGrande, PIQA and ARC-Easy—which is the signature of noise rather than of a pattern.

At a 5 B-token budget the downstream suite sits near its noise floor: WinoGrande scores 50.0 and 51.6 against a 50.0 chance floor, so roughly two of five tasks measure nothing. Validation bits-per-byte, computed over 645 M held-out tokens, separates architectures far more sensitively at this scale, which is why it was pre-registered.

The honest summary is that **the hybrid matches on quality and wins decisively on decode**. It is not a downstream-task win, and is not presented as one.

B. Per-domain difficulty

Table VI decomposes validation bits-per-byte by source for the dense arm. The spread is nearly a factor of two and is informative for anyone reusing this corpus: code and encyclopedic text are far more predictable than broad filtered web text, so an aggregate figure is dominated by the mixture weights as much as by the model.

C. CPU decode speed

Decode is measured with 4-bit weights on 8 threads, generating 128 tokens after priming a context of the stated depth. Arms alternate within a single pass so that background load on the measurement machine perturbs both equally; the ratio within a pass is therefore more trustworthy than the absolute throughputs, which are depressed by whatever else is running.

Table VII contains the central result, and its *shape* is the thesis. At depth zero the hybrid has almost nothing to gain, because its advantage *is* the cache it does not keep and an empty context has no cache to re-read. The advantage grows monotonically with context, reaching 1.76× at the trained context length.

One confound runs against the hybrid: the dense twin receives the reference runtime’s better-optimised attention graph, which has had far more engineering attention than the hybrid’s convolution path. The margin is achieved in spite of that, and

TABLE VII
CPU DECODE, TRAINED WEIGHTS, ALTERNATING PASSES. THE ADVANTAGE GROWS MONOTONICALLY WITH CONTEXT.

Depth	Hybrid	Dense twin	Ratio
0 (empty)	1111.9 ± 25.7	922.8 ± 14.2	1.20×
512	960.3 ± 10.7	664.4 ± 6.5	1.45×
2048	739.3 ± 35.5	420.3 ± 5.5	1.76×




Fig. 2. Decode throughput against context depth (4-bit weights, 8 threads, trained weights). Both curves fall as the context grows, but the dense twin falls roughly twice as fast: it re-reads $\kappa = 12,288$ B of cache per context token against the hybrid’s 6,144 B. The vertical gap, not the absolute height, is the architectural result.

would be expected to widen if the convolution kernels received comparable optimisation.

Two figures appearing in this project’s history should not be quoted: 1.29× originated from a non-alternating measurement that does not reproduce, and 1.15–1.17× is the depth-zero row, a floor rather than a result.

Fig. 2 plots the same data. The two curves do not merely differ by a constant: they diverge, and the divergence is the quantity the architecture was designed to produce.

D. Comparison against an external peer

Daedalus decodes 2.08× faster at the context it is built for while carrying 19 % more parameters. The same signature—near unity at depth zero, growing with context—reproduces against a different architecture written by a different author, which is stronger evidence for the mechanism than the within-project comparison alone, since it cannot be an artefact of this project’s training or export code.

E. Quantised artefact size

Four-bit output size depends on tensor shapes rather than weight values, so this result was final before training completed. Matched within 0.5 % at half precision (306.22 MiB against 307.63 MiB), the hybrid’s shipped file is 95.56 MiB against the twin’s 101.62 MiB—**6.3 % smaller**, at 4.99 versus 5.29 bits per weight. The quantiser reports that one of 266 tensors required fallback quantisation on the dense twin and issues no such warning on the hybrid: a shape effect that happens to favour the hybrid’s dimensions, and a reminder that

TABLE VIII
CPU DECODE AGAINST AN EXTERNAL 135 M-PARAMETER PEER, SAME HARNESS.

Depth	Daedalus	Peer-135M	Ratio
0	960.9 $\pm$ 3.4	908.1 $\pm$ 37.1	1.06 $\times$
512	933.7 $\pm$ 28.5	625.7 $\pm$ 38.5	1.49 $\times$
2048	648.6 $\pm$ 12.6	312.4 $\pm$ 7.2	2.08$\times$

TABLE IX
FIVE-TASK MEAN. EVERY PEER WAS RE-SCORED HERE RATHER THAN QUOTED FROM ITS OWN PAPER, SO THE COLUMN IS INTERNALLY COMPARABLE.

Model	Training tokens	5-task
Daedalus-150M	59.9 B	47.31
MobileLLM-125M	1 T	46.3 (published)
GPT-2 124M	—	42.2
OPT-125M	180 B	42.1
GPT-neo-125M	300 B	41.9
Pythia-160M	300 B	41.0
Peer-135M	2 T	51.2

the deployed artefact size is not a pure function of parameter count.

F. One datapoint against the bar

The hybrid ablation arm—identical architecture, fully decayed, but trained on only 5 B tokens—scores **44.7** on the five-task harness. That clears the 42.2 bar and exceeds all four 300 B-token-class peers on every one of the five tasks. This is not the Daedalus result and nothing is projected from it; it moves the open question from whether the architecture can reach the bar to how far beyond it twelve times more data carries.

VIII. HEADLINE RESULTS

Table IX gives the finished model measured on the harness described in Section XI, alongside the peers it was required to beat.

The model clears the 42.20 bar by 5.11 points and beats every model it was set against, each of which saw three to six times more data. It also exceeds MobileLLM-125M’s published figure, which came from a trillion tokens.

The 2 T-token peer in the last row remains ahead by 3.9 points. That was conceded before training started: this project trades quality against decode speed at a fixed size, and Section VII-C is the other half of that trade.

Validation bits-per-byte over a 645 M-token held-out set is **0.8685**. The same architecture trained on 5 B tokens reached 0.9104, so the remaining 55 B tokens bought a 4.6 % improvement — a large gain by the standards of this metric, which moves slowly.

Quantising to 4 bits costs roughly 6 % perplexity (9.18 at half precision against 9.75 at Q4_0, measured on held-out encyclopedic text). That is higher than the 2.5 % measured at 5 B tokens and is discussed in Section IX-C.

TABLE X
PER-TASK SCORES. WINOGRANDE SITS AT ITS CHANCE FLOOR AND IS MEASURING NOTHING AT THIS SCALE.

Task	Score	Chance
PIQA	65.78	50.0
ARC-Easy	50.42	25.0
WinoGrande	50.04	50.0
HellaSwag	37.93	25.0
OpenBookQA	32.40	25.0
Mean	47.31	

IX. LIMITATIONS

A. The training mixture drifted from its target

The corpus holds 16.9 B unique tokens and the run consumed 59.9 B, so most text is seen about three and a half times. Each source is capped at four repetitions, and for the largest sources that cap binds — their share cannot be met without exceeding it, so the freed share flows to smaller sources that still have unread text.

The realised mixture therefore sits some distance from the intended one. Measured as L_1 distance in percentage points, the run as a whole is at **10.42** against a limit of 10.0 that was fixed in advance. The largest single-source deviation is -0.24 points, which is negligible against benchmark noise, but the headline figure is over the limit and is reported as such. If downstream results ever look weaker than the peer comparison suggests they should, mixture drift is a live explanation.

B. The last 8 % of training used a slightly smaller corpus

Training was interrupted and resumed from a checkpoint. Three consequences are permanent properties of the released weights and are stated rather than smoothed over.

The optimiser state was not carried across: momentum restarted from zero and re-warmed over a few hundred steps. This is mild — by that point the learning rate had already decayed to 0.0035 — but the run is not one continuous optimiser trajectory.

The data cursor also reset, so the final 4.8 B tokens are a fresh pass over the corpus rather than a continuation. Some documents were seen twice and some not at all.

Those tokens came from a corpus snapshot 0.42 B tokens smaller than the one used earlier, which is what pushed the mixture figure above from 9.94 to 10.42.

C. Quantisation-aware training did not run

The plan was to spend the final 5 % of training with fake quantisation applied on the deployment format’s exact 4-bit grid, verified against the runtime’s own kernel rather than reimplemented. It produced a non-finite loss on its first step and was disabled; the run finished without it.

The released model is therefore quantised after training rather than during it, and carries the full 4-bit penalty — about 6 % perplexity rather than the 2.5 % measured at smaller scale. The half-precision weights are published, so this is fixable without repeating any training: either by running the

quantisation-aware pass from the finished checkpoint, or by choosing a format with better error behaviour at the same bit width. We did not diagnose the failure.

D. The vocabulary is larger than this size warrants

The tokeniser, and with it a 49,152-entry vocabulary, was adopted from an existing model because an earlier plan required student and teacher to share one exactly. That plan was dropped for unrelated reasons; the vocabulary stayed.

Scaling laws put the optimum for a 150 M model nearer 24–32 k. At 49,152 the embedding table holds 37.7 M parameters — **23 % of the model** — where a 32 k vocabulary would hold 24.6 M, freeing roughly 13 M parameters for layers that compute rather than look up. Tied embeddings recover half the waste, which is probably why it was never revisited. A smaller vocabulary is the first thing to change in a successor.

We did verify the tokeniser rather than assume it: reading token ids directly out of the quantised file gave **zero mismatches** against the reference tokeniser. A silent mismatch here ships a model that produces fluent nonsense while every reported metric looks healthy, and nothing else in the pipeline would catch it.

E. Single seed, English only

Every number here comes from one seed. The 0.81 % ablation margin is not a confidence interval, and the downstream comparison is reported as a tie partly because one seed cannot resolve a quarter of a standard deviation. The model is English-only, and the decode advantage is measured at the 2048-token context it was trained for, not extrapolated past it.

X. A NEGATIVE RESULT: DEAD CHANNELS CANNOT BE RECLAIMED

Approximately 47.9 % of short-convolution channels contribute nothing to the model’s output. This is a stable plateau rather than progressive decay: the fraction measures 47.928 % at step 9,896 and 47.993 % at step 30,041, and is unchanged over the following 10,743 steps. It represents roughly 13.6 M inert parameters, an 8.5 % parameter inefficiency.

Because the mask is stable and the convolution is depthwise, structural pruning at export appeared straightforward: narrow each block’s projections and kernel while leaving the residual stream at full width, for an estimated 7.7 MB saving. **This was tested and does not work.** The reference runtime creates all three short-convolution tensors at fixed model width and shape-checks them at load time; a real artefact narrowed from 768 to 640 is rejected with `check_tensor_dims: expected 3,768, got 3,640`, while both the unmodified file and the same file rebuilt at full width through the identical writer load correctly. That second control is what makes the finding conclusive: the rejection is the narrowing itself, not an artefact of the rewriting tool.

Reclaiming the parameters would require patching the inference runtime, which would forfeit stock-binary compatibility—the foundation of the CPU-decode claim—in exchange for 7.7 MB. Four-bit quantisation moreover spends

four bits on a zero exactly as on any other weight, so the inert channels cost nothing beyond their share of the file. The appropriate remedy belongs to the next model’s initialisation and regularisation, not to a retrofit of this one.

XI. EVALUATION METHODOLOGY

Because several claims in this paper are comparative, the instruments deserve description.

A. Quality

The five-task suite is scored under lm-evaluation-harness conventions with length-normalised accuracy where the harness specifies it. Crucially, *peer models are re-scored on the same harness* rather than compared against published figures. The gap between the two is not negligible: published eight-task means for the peer set run 0.5–1.5 points above what the same checkpoints score here, because the task subset and normalisation differ. Comparing a locally measured number against a published one would manufacture roughly a point of spurious advantage.

Validation bits-per-byte is computed over a 645 M-token held-out set. Bits per *byte*, rather than per token, is used so that the figure is comparable across tokenisers—a model with a more efficient tokeniser would otherwise appear better at equal predictive quality.

B. Decode

Decode throughput is measured by generating 128 tokens after priming a context of the stated depth, on 8 threads, with 4-bit weights, repeated and reported as mean $\pm$ standard deviation. Two protocol details matter. First, arms *alternate within a single pass* rather than running back to back: an earlier back-to-back measurement in this project produced a $1.29\times$ figure that did not reproduce, because background load drifted between the two halves. Second, absolute throughputs are depressed by whatever else shares the machine, so the ratio within a pass is the trustworthy quantity and the absolute numbers should be treated as a floor.

The measurement includes only generation, not prompt processing. Prompt processing is compute-bound and parallel over positions, so it is a regime where attention is not disadvantaged; including it would flatter the hybrid on workloads that decode little.

XII. DEPLOYMENT NOTES

The model targets stock inference binaries with no patches, which constrained several decisions and is worth making explicit for anyone reproducing the deployment.

Quantisation format. Q4_0 is chosen over formats with better error characteristics at equal bit width because its dot-product kernels are the best optimised on the target hardware. The measured quality cost of that choice is 2.5 % perplexity at 5 B tokens without quantisation-aware training; the throughput it buys is what the entire design exists to deliver. Formats such as Q4_K would reduce the error and give back part of the speed.

Memory footprint. The shipped artefact is 95.56 MiB. At 2048-token context the key–value cache adds approximately 12.6 MB, so a complete single-user session fits comfortably under 128 MB—small enough for the model to be resident alongside an application rather than requiring a dedicated process boundary.

Thread scaling. All figures here use 8 threads. Because decoding at batch size one is memory-bound, throughput saturates once threads suffice to saturate memory bandwidth; adding cores past that point yields little. The hybrid saturates at a lower thread count than the dense twin, since it moves fewer bytes per token, which is a secondary benefit on shared machines.

XIII. DISCUSSION

When the hybrid wins. The advantage is a function of context depth. Any deployment that keeps a long conversation, retrieves documents into the prompt, or processes files will sit at the right-hand end of Table VII. Deployments that issue short, independent prompts into an empty context will see the depth-0 row and little benefit.

Extrapolation beyond 2048. Eq. (4) implies the ratio keeps growing with t , and the measured trend is consistent with that. We decline to quote a figure beyond the trained context, because the model has not been trained there and the claim would not be about a usable model.

Where the ratio could be improved. Two thirds of the layers are already cache-free, so the remaining cache is concentrated in six layers. Reducing attention further trades against retrieval capability, which this experiment did not measure; a retrieval-sensitive evaluation is the right instrument for that question and is future work.

Future work. Diagnosing the quantisation-aware training failure; addressing dead channels at initialisation rather than export; a depth ablation (18×768 against 24×640) which is designed but unrun; multi-seed replication of the ablation; and a retrieval-sensitive evaluation to bound how far the attention fraction can fall.

XIV. REPRODUCIBILITY

Every number in this paper comes from a file in the public repository. The head-to-head quality, decode and file-size figures come from the comparison’s own results file; the winning condition comes from a document timestamped before either model was scored, so it could not have been fitted to the outcome; the five-task and bits-per-byte figures come from the evaluation outputs; and the pruning result comes from a script that can be re-run against a real model file.

The model weights are published at half precision as well as 4-bit, so the quantisation figures can be reproduced and improved on without repeating any training. The corpus is assembled entirely from public datasets.

XV. CONCLUSION

The proposition was specific and falsifiable: for one user decoding on a CPU, replacing two thirds of a transformer’s

attention layers with fixed-state short convolutions buys a large speed advantage at no cost in quality.

Both halves held. The parameter-matched comparison gave the hybrid the chosen quality metric by 0.81 % against a margin fixed in advance, a tie on downstream tasks, a 6.3 % smaller 4-bit file, and **1.76 $\times$ faster decoding** at 2048 tokens — 2.08 $\times$ against an outside model. The advantage grows with context in every measurement, which is what identifies the cause: a model that were simply leaner would be faster by a constant.

At full scale the model scores **47.31** against a bar of 42.20 set before training, beating every model in its size class trained on three to six times more data. Bits-per-byte is 0.8685.

What remains open is not the architecture but the engineering around it: a 4-bit penalty that quantisation-aware training was meant to remove, roughly half the convolution channels sitting inert because nothing in training discourages that, and a vocabulary inherited rather than chosen. None of the three requires new research to fix, and all three are cheaper to address in the next model than to retrofit into this one.

REFERENCES

- [1] K. Bergsma *et al.*, “Straight to zero: Why linearly decaying the learning rate to zero works best for LLMs,” *arXiv:2502.15938*, 2025.
- [2] N. Muennighoff *et al.*, “Scaling data-constrained language models,” *Advances in Neural Information Processing Systems*, 2023, arXiv:2305.16264.
- [3] J. Ainslie *et al.*, “GQA: Training generalized multi-query transformer models from multi-head checkpoints,” *Proc. EMNLP*, 2023, arXiv:2305.13245.
- [4] K. Jordan *et al.*, “Muon: An optimizer for hidden layers in neural networks,” 2024.
- [5] G. Gerganov *et al.*, “llama.cpp: LLM inference in C/C++,” 2023. [Online]. Available: <https://github.com/ggml-org/llama.cpp>
- [6] S. Biderman *et al.*, “Pythia: A suite for analyzing large language models across training and scaling,” *Proc. ICML*, 2023, arXiv:2304.01373.
- [7] Z. Liu *et al.*, “MobileLLM: Optimizing sub-billion parameter language models for on-device use cases,” *Proc. ICML*, 2024, arXiv:2402.14905.
- [8] A. Gu and T. Dao, “Mamba: Linear-time sequence modeling with selective state spaces,” *arXiv:2312.00752*, 2023.
- [9] B. Peng *et al.*, “RWKV: Reinventing RNNs for the transformer era,” *Findings of EMNLP*, 2023, arXiv:2305.13048.
- [10] S. De *et al.*, “Griffin: Mixing gated linear recurrences with local attention for efficient language models,” *arXiv:2402.19427*, 2024.